

# Deep Learning

## Natural Language Processing with Deep Learning


Alex Olson

Adapted from material by Charles Ollion & Olivier Grisel

# Natural Language Processing



[Google Translate System - 2016]



[Socher 2015]

# Natural Language Processing

- Sentence/Document level Classification (topic, sentiment)

# Natural Language Processing

- Sentence/Document level Classification (topic, sentiment)
- Topic modeling (LDA, ...)

# Natural Language Processing

- Sentence/Document level Classification (topic, sentiment)
- Topic modeling (LDA, ...)
- Translation

# Natural Language Processing

- Sentence/Document level Classification (topic, sentiment)
- Topic modeling (LDA, ...)
- Translation
- Chatbots / dialogue systems / assistants (Alexa, ...)

# Natural Language Processing

- Sentence/Document level Classification (topic, sentiment)
- Topic modeling (LDA, ...)
- Translation
- Chatbots / dialogue systems / assistants (Alexa, ...)
- Summarization

# Natural Language Processing

- Sentence/Document level Classification (topic, sentiment)
- Topic modeling (LDA, ...)
- Translation
- Chatbots / dialogue systems / assistants (Alexa, ...)
- Summarization

Useful open source projects





# Outline

Classification and word representation

# Outline

Classification and word representation

Word2Vec

# Outline

Classification and word representation

Word2Vec

Language Modelling

# Outline

Classification and word representation

Word2Vec

Language Modelling

Recurrent neural networks

# Word Representation and Word2Vec

# Word representation

Words are indexed and represented as 1-hot vectors

# Word representation

Words are indexed and represented as 1-hot vectors

Large Vocabulary of possible words  $|V|$

# Word representation

Words are indexed and represented as 1-hot vectors

Large Vocabulary of possible words  $|V|$

Use of **Embeddings** as inputs in all Deep NLP tasks



# Word representation


Words are indexed and represented as 1-hot vectors

Large Vocabulary of possible words  $|V|$

Use of **Embeddings** as inputs in all Deep NLP tasks


Word embeddings usually have dimensions 50, 100, 200, 300

# Supervised Text Classification



Joulin, Armand, et al. "Bag of tricks for efficient text classification." FAIR 2016

# Supervised Text Classification




$E$  embedding (linear projection)

$|V| \times H$

Joulin, Armand, et al. "Bag of tricks for efficient text classification." FAIR 2016

# Supervised Text Classification



$E$  embedding (linear projection)


$|V| \times H$

Embeddings are averaged

hidden activation size:  $H$

Joulin, Armand, et al. "Bag of tricks for efficient text classification." FAIR 2016

# Supervised Text Classification



$E$  embedding (linear projection)

$|V| \times H$

Embeddings are averaged


hidden activation size:  $H$

Dense output connection  $W, b$

$H \times K$

Joulin, Armand, et al. "Bag of tricks for efficient text classification." FAIR 2016

# Supervised Text Classification



**E** embedding (linear projection)

$|V| \times H$

Embeddings are averaged

hidden activation size: **H**


Dense output connection **W**, **b**

**H**  $\times$  **K**

Softmax and cross-entropy loss

Joulin, Armand, et al. "Bag of tricks for efficient text classification." FAIR 2016


# Supervised Text Classification



- Very efficient (speed and accuracy) on large datasets

Joulin, Armand, et al. "Bag of tricks for efficient text classification." FAIR 2016

# Supervised Text Classification




- Very efficient (speed and accuracy) on large datasets
- State-of-the-art (or close to) on several classification, when adding bigrams/trigrams

Joulin, Armand, et al. "Bag of tricks for efficient text classification." FAIR 2016



# Supervised Text Classification



- Very efficient (speed and accuracy) on large datasets
- State-of-the-art (or close to) on several classification, when adding bigrams/trigrams
- Little gains from depth

Joulin, Armand, et al. "Bag of tricks for efficient text classification." FAIR 2016

# Transfer Learning for Text

Similar to image: can we have word representations that are generic enough to transfer from one task to another?

# Transfer Learning for Text

Similar to image: can we have word representations that are generic enough to transfer from one task to another?

Unsupervised / self-supervised learning of word representations

# Transfer Learning for Text

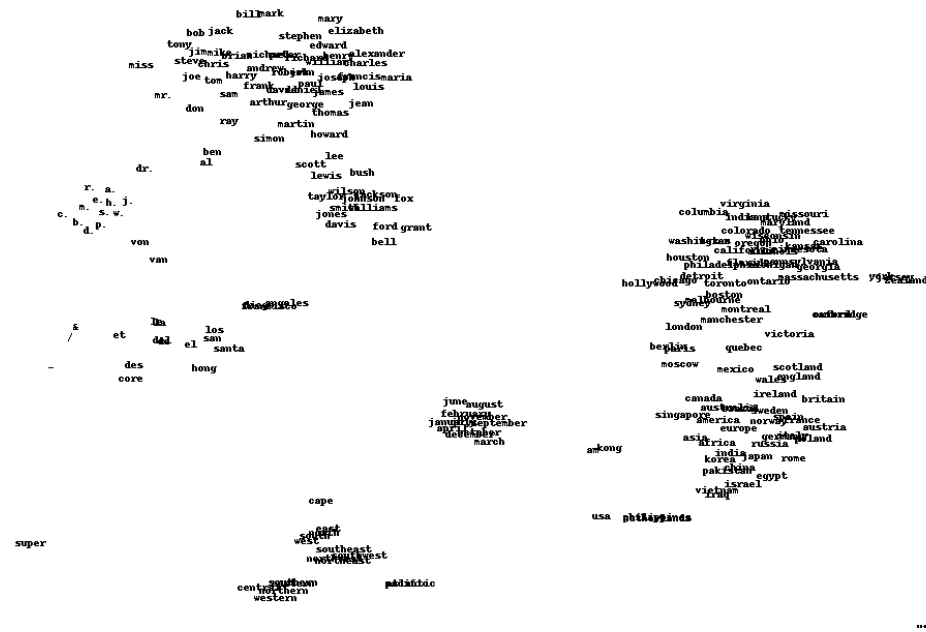
Similar to image: can we have word representations that are generic enough to transfer from one task to another?

Unsupervised / self-supervised learning of word representations

Unlabelled text data is almost infinite:

- Wikipedia dumps
- Project Gutenberg
- Social Networks
- Common Crawl

# Word Vectors



excerpt from work by J. Turian on a model trained by R. Collobert et al. 2008

# Word2Vec

FRANCE	JESUS	XBOX	REDDISH	SCRATCHED	MEGABITS
AUSTRIA	GOD	AMIGA	GREENISH	NAILED	OCTETS
BELGIUM	SATI	PLAYSTATION	BLUISH	SMASHED	MB/s
GERMANY	CHRIST	MSX	PINKISH	PUNCHED	BIT/s
ITALY	SATAN	IPOD	PURPLISH	POPPED	BAUD
GREECE	KALI	SEGA	BROWNISH	CRIMPED	CARATS
SWEDEN	INDRA	PSNUMBER	GREYISH	SCRAPED	KBIT/s
NORWAY	VISHNU	HD	GRAYISH	SCREWED	MEGAHERTZ
EUROPE	ANANDA	DREAMCAST	WHITISH	SECTIONED	MEGAPIXELS
HUNGARY	PARVATI	GEFORCE	SILVERY	SLASHED	GBIT/s
SWITZERLAND	GRACE	CAPCOM	YELLOWISH	RIPPED	AMPERES

Colobert et al. 2011, Mikolov, et al. 2013

# Word2Vec


FRANCE	JESUS	XBOX	REDDISH	SCRATCHED	MEGABITS
AUSTRIA	GOD	AMIGA	GREENISH	NAILED	OCTETS
BELGIUM	SATI	PLAYSTATION	BLUISH	SMASHED	MB/s
GERMANY	CHRIST	MSX	PINKISH	PUNCHED	BIT/s
ITALY	SATAN	IPOD	PURPLISH	POPPED	BAUD
GREECE	KALI	SEGA	BROWNISH	CRIMPED	CARATS
SWEDEN	INDRA	PSNUMBER	GREYISH	SCRAPED	KBIT/s
NORWAY	VISHNU	HD	GRAYISH	SCREWED	MEGAHERTZ
EUROPE	ANANDA	DREAMCAST	WHITISH	SECTIONED	MEGAPIXELS
HUNGARY	PARVATI	GEFORCE	SILVERY	SLASHED	GBIT/s
SWITZERLAND	GRACE	CAPCOM	YELLOWISH	RIPPED	AMPERES

## Compositionality

Czech + currency	Vietnam + capital	German + airlines	Russian + river	French + actress
koruna	Hanoi	airline Lufthansa	Moscow	Juliette Binoche
Check crown	Ho Chi Minh City	carrier Lufthansa	Volga River	Vanessa Paradis
Polish zolty	Viet Nam	flag carrier Lufthansa	upriver	Charlotte Gainsbourg
CTK	Vietnamese	Lufthansa	Russia	Cecile De

Colobert et al. 2011, Mikolov, et al. 2013


# Word Analogies



Mikolov, Tomas, et al. "Distributed representations of words and phrases and their compositionality." NIPS 2013




# Word Analogies



- Linear relations in Word2Vec embeddings

Mikolov, Tomas, et al. "Distributed representations of words and phrases and their compositionality." NIPS 2013

# Word Analogies



- Linear relations in Word2Vec embeddings
- Many come from text structure (e.g. Wikipedia)

Mikolov, Tomas, et al. "Distributed representations of words and phrases and their compositionality." NIPS 2013

# Self-supervised training

Distributional Hypothesis (Harris, 1954): *“words are characterised by the company that they keep”*

Main idea: learning word embeddings by predicting word contexts

# Self-supervised training

Distributional Hypothesis (Harris, 1954): *“words are characterised by the company that they keep”*

Main idea: learning word embeddings by predicting word contexts

Given a word e.g. “carrot” and any other word  $w \in V$  predict probability  $P(w|\text{carrot})$  that  $w$  occurs in the context of “carrot”.

# Self-supervised training

Distributional Hypothesis (Harris, 1954): *“words are characterised by the company that they keep”*

Main idea: learning word embeddings by predicting word contexts

Given a word e.g. “carrot” and any other word  $w \in V$  predict probability  $P(w|\text{carrot})$  that  $w$  occurs in the context of “carrot”.

- Unsupervised / self-supervised: no need for class labels.
- (Self-)supervision comes from context.
- Requires a lot of text data to cover rare words correctly.

# Word2Vec: CBoW

CBoW: representing the context as Continuous Bag-of-Words

Self-supervision from large unlabeled corpus of text: *slide* over an anchor word and its context:

the carrot is a root vegetable, usually orange


Mikolov, Tomas, et al. "Distributed representations of words and phrases and their compositionality." NIPS 2013

# Word2Vec: CBoW

CBoW: representing the context as Continuous Bag-of-Words


Self-supervision from large unlabeled corpus of text: *slide* over an anchor word and its context:

the carrot is a root vegetable, usually orange



Mikolov, Tomas, et al. "Distributed representations of words and phrases and their compositionality." NIPS 2013


# Word2Vec: Skip Gram



- Given the central word, predict occurrence of other words in its context.



# Word2Vec: Skip Gram



- Given the central word, predict occurrence of other words in its context.
- Widely used in practice

# Word2Vec: Negative Sampling

- Task is simplified further: binary classification of word pairs

# Word2Vec: Negative Sampling

- Task is simplified further: **binary classification** of word pairs
- For the sentence "The quick brown fox jumps over the lazy dog":

# Word2Vec: Negative Sampling

- Task is simplified further: **binary classification** of word pairs
- For the sentence "The quick brown fox jumps over the lazy dog":
- "quick" and "fox" are positive examples (if context window is 2)
- "quick" and "apple" are negative examples

# Word2Vec: Negative Sampling

- Task is simplified further: **binary classification** of word pairs
- For the sentence "The quick brown fox jumps over the lazy dog":
- "quick" and "fox" are positive examples (if context window is 2)
- "quick" and "apple" are negative examples
- By sampling negative examples, we don't just bring similar words' embeddings closer, but also push away dissimilar words' embeddings.

# Transformer-based methods

- Attention mechanism: more recent and more powerful than Word2Vec
- BERT (Bidirectional Encoder Representations from Transformers) allows for contextual embeddings (different embeddings for the same word in different contexts)
- For example, "bank" in "river bank" and "bank account" will have different embeddings
- This means converting a word to a vector is no longer a simple lookup in a table, but a function of the entire sentence

# Transformer-based methods

- Sub-word tokenization: BERT uses a sub-word tokenization, which allows it to handle out-of-vocabulary words better than Word2Vec
- For example, "unbelievable" can be split into "un" and "believable"
- This means that the model can guess the meaning of words it has never seen before, based on the meanings of their parts
- OpenAI tokenization example: <https://platform.openai.com/tokenizer>

# Take Away on Embeddings

For text applications, inputs of Neural Networks are Embeddings



# Take Away on Embeddings

For text applications, inputs of Neural Networks are Embeddings

- If little training data and a wide vocabulary not well covered by training data, use pre-trained self-supervised embeddings (word2vec, or with more time and resources, BERT, GPT, etc.)

# Take Away on Embeddings

For text applications, inputs of Neural Networks are Embeddings

- If little training data and a wide vocabulary not well covered by training data, use pre-trained self-supervised embeddings (word2vec, or with more time and resources, BERT, GPT, etc.)
- If large training data with labels, directly learn task-specific embedding for more precise representation.

# Take Away on Embeddings

For text applications, inputs of Neural Networks are Embeddings

- If little training data and a wide vocabulary not well covered by training data, use pre-trained self-supervised embeddings (word2vec, or with more time and resources, BERT, GPT, etc.)
- If large training data with labels, directly learn task-specific embedding for more precise representation.
- word2vec uses Bag-of-Words (BoW): they ignore the order in word sequences

# Take Away on Embeddings

For text applications, inputs of Neural Networks are Embeddings

- If little training data and a wide vocabulary not well covered by training data, use pre-trained self-supervised embeddings (word2vec, or with more time and resources, BERT, GPT, etc.)
- If large training data with labels, directly learn task-specific embedding for more precise representation.
- word2vec uses Bag-of-Words (BoW): they ignore the order in word sequences
- Depth & non-linear activations on hidden layers are not that useful for BoW text classification.

# Language Modelling and Recurrent Neural Networks

# Language Models

Assign a probability to a sequence of words, such that plausible sequences have higher probabilities e.g:

- $p(\text{"I like cats"}) > p(\text{"I table cats"})$
- $p(\text{"I like cats"}) > p(\text{"like I cats"})$

# Language Models

Assign a probability to a sequence of words, such that plausible sequences have higher probabilities e.g:

- $p(\text{"I like cats"}) > p(\text{"I table cats"})$
- $p(\text{"I like cats"}) > p(\text{"like I cats"})$

Likelihoods are factorized:

$$p_{\theta}(w_0)$$

$p_{\theta}$  is parametrized by a neural network.

# Language Models

Assign a probability to a sequence of words, such that plausible sequences have higher probabilities e.g:

- $p(\text{"I like cats"}) > p(\text{"I table cats"})$
- $p(\text{"I like cats"}) > p(\text{"like I cats"})$

Likelihoods are factorized:

$$p_{\theta}(w_0) \cdot p_{\theta}(w_1 | w_0)$$

$p_{\theta}$  is parametrized by a neural network.



# Language Models

Assign a probability to a sequence of words, such that plausible sequences have higher probabilities e.g:

- $p(\text{"I like cats"}) > p(\text{"I table cats"})$
- $p(\text{"I like cats"}) > p(\text{"like I cats"})$

Likelihoods are factorized:

$$p_{\theta}(w_0) \cdot p_{\theta}(w_1|w_0) \cdot \dots \cdot p_{\theta}(w_n|w_{n-1}, w_{n-2}, \dots, w_0)$$

$p_{\theta}$  is parametrized by a neural network.

# Language Models

Assign a probability to a sequence of words, such that plausible sequences have higher probabilities e.g:

- $p(\text{"I like cats"}) > p(\text{"I table cats"})$
- $p(\text{"I like cats"}) > p(\text{"like I cats"})$

Likelihoods are factorized:

$$p_{\theta}(w_0) \cdot p_{\theta}(w_1 | w_0) \cdot \dots \cdot p_{\theta}(w_n | w_{n-1}, w_{n-2}, \dots, w_0)$$

$p_{\theta}$  is parametrized by a neural network.

The internal representation of the model can better capture the meaning of a sequence than a simple Bag-of-Words.

# Conditional Language Models

NLP problems expressed as Conditional Language Models:

Translation:  $p(\textit{Target}|\textit{Source})$

- *Source*: "J'aime les chats"
- *Target*: "I like cats"

# Conditional Language Models

NLP problems expressed as Conditional Language Models:

Translation:  $p(\textit{Target}|\textit{Source})$

- *Source*: "J'aime les chats"
- *Target*: "I like cats"

Model the output word by word:

$p_{\theta}(w_0|\textit{Source})$

# Conditional Language Models

NLP problems expressed as Conditional Language Models:

Translation:  $p(\textit{Target}|\textit{Source})$

- *Source*: "J'aime les chats"
- *Target*: "I like cats"

Model the output word by word:

$$p_{\theta}(w_0|\textit{Source}) \cdot p_{\theta}(w_1|w_0, \textit{Source}) \cdot \dots$$

# Conditional Language Models

Question Answering / Dialogue:

$p(\text{Answer} | \text{Question}, \text{Context})$

- *Context:*
  - "John puts two glasses on the table."
  - "Bob adds two more glasses."
  - "Bob leaves the kitchen to play baseball in the garden."
- *Question:* "How many glasses are there?"
- *Answer:* "There are four glasses."

# Conditional Language Models

Question Answering / Dialogue:


$p(\text{Answer} | \text{Question}, \text{Context})$

- *Context:*
  - "John puts two glasses on the table."
  - "Bob adds two more glasses."
  - "Bob leaves the kitchen to play baseball in the garden."
- *Question:* "How many glasses are there?"
- *Answer:* "There are four glasses."

Image Captionning:  $p(\text{Caption} | \text{Image})$


- Image is usually the 2048-d representation from a CNN

# Simple Language Model






# Simple Language Model



Fixed context size

- **Average embeddings:** (same as CBoW) no sequence information


# Simple Language Model



Fixed context size

- **Average embeddings:** (same as CBoW) no sequence information
- **Concatenate embeddings:** introduces many parameters


# Simple Language Model




Fixed context size

- **Average embeddings:** (same as CBoW) no sequence information
- **Concatenate embeddings:** introduces many parameters
- Still does not take well into account varying sequence sizes and sequence dependencies


# Recurrent Neural Network




# Recurrent Neural Network




Unroll over a sequence  $(x_0, x_1, x_2)$ :




# Recurrent Neural Network




Unroll over a sequence  $(x_0, x_1, x_2)$ :



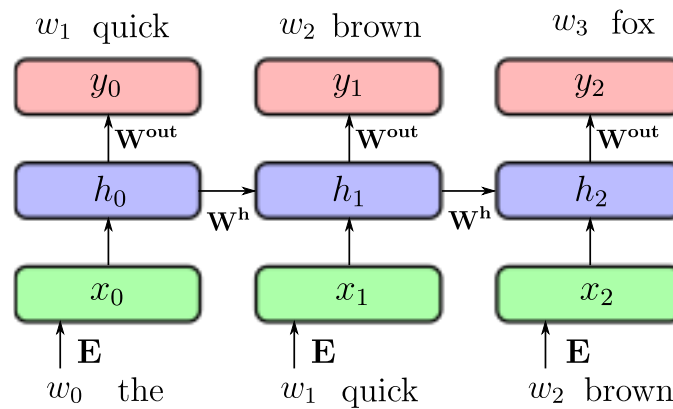
# Recurrent Neural Network



Unroll over a sequence  $(x_0, x_1, x_2)$ :



# Language Modelling

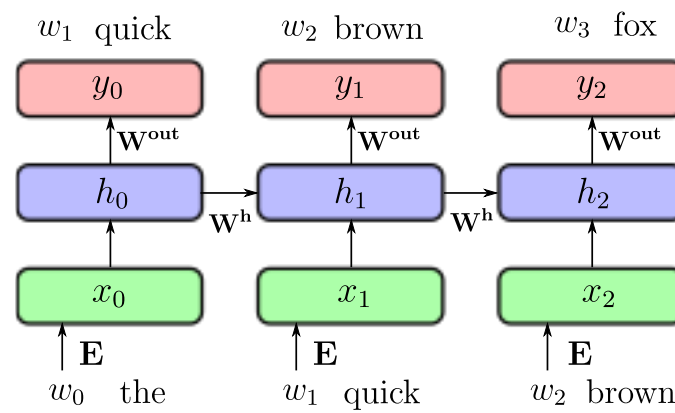


**input** ( $w_0, w_1, \dots, w_t$ ) sequence of words ( 1-hot encoded )

**output** ( $w_1, w_2, \dots, w_{t+1}$ ) shifted sequence of words ( 1-hot encoded )



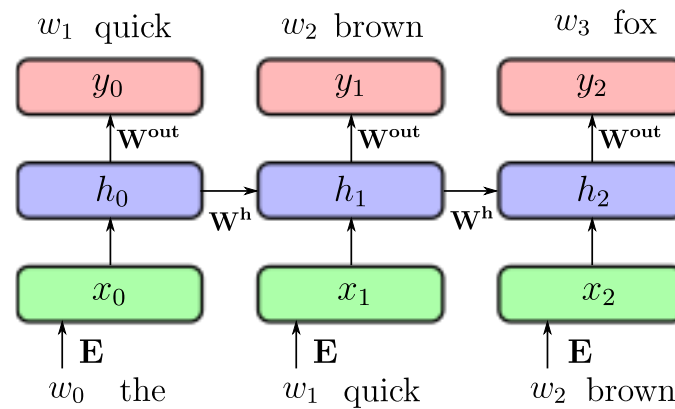
# Language Modelling



$$x_t = \text{Emb}(w_t) = \mathbf{E}w_t$$

input projection  $\mathbf{H}$

# Language Modelling




$$x_t = \text{Emb}(w_t) = \mathbf{E}w_t$$

$$h_t = g(\mathbf{W}^h h_{t-1} + x_t + b^h)$$

input projection  $\mathbf{H}$

recurrent connection  $\mathbf{H}$

# Language Modelling



$$x_t = \text{Emb}(w_t) = \mathbf{E}w_t$$

$$h_t = g(\mathbf{W}^h h_{t-1} + x_t + b^h)$$

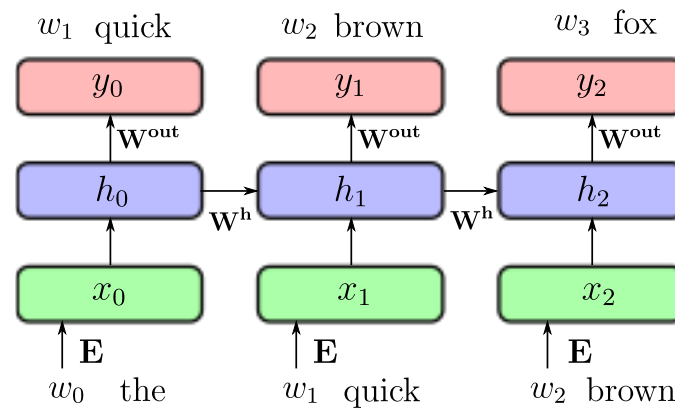
$$y = \text{softmax}(\mathbf{W}^o h_t + b^o)$$

input projection  $\mathbf{H}$

recurrent connection  $\mathbf{H}$

output projection  $\mathbf{K} = |\mathbf{V}|$

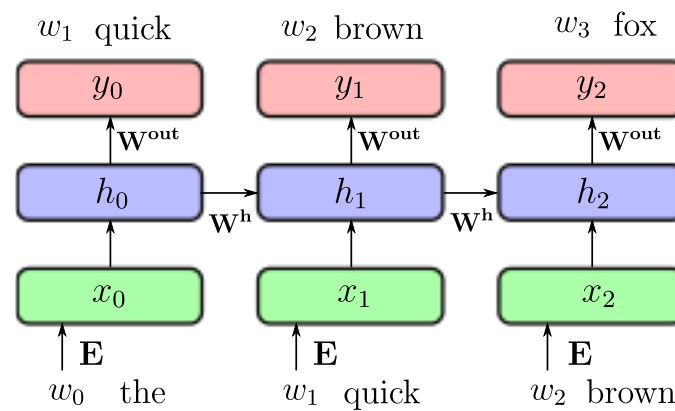
# Recurrent Neural Network



Input embedding  $\mathbf{E}$

$|\mathbf{V}| \times \mathbf{H}$

# Recurrent Neural Network




Input embedding  $\mathbf{E}$

$|\mathbf{V}| \times \mathbf{H}$

Recurrent weights  $\mathbf{W}^h$

$\mathbf{H} \times \mathbf{H}$

# Recurrent Neural Network



Input embedding  $\mathbf{E}$

$|V| \times H$

Recurrent weights  $\mathbf{W}^h$


$H \times H$

Output weights  $\mathbf{W}^{\text{out}}$

$H \times K = H \times |V|$


# Backpropagation through time

Similar as standard backpropagation on unrolled network



# Backpropagation through time


Similar as standard backpropagation on unrolled network






# Backpropagation through time

Similar as standard backpropagation on unrolled network



# Backpropagation through time


Similar as standard backpropagation on unrolled network



- Similar as training **very deep networks** with tied parameters
- Example between  $x_0$  and  $y_2$ :  $W^h$  is used twice

# Backpropagation through time


Similar as standard backpropagation on unrolled network



- Similar as training **very deep networks** with tied parameters
- Example between  $x_0$  and  $y_2$ :  $W^h$  is used twice
- Usually truncate the backprop after  $T$  timesteps

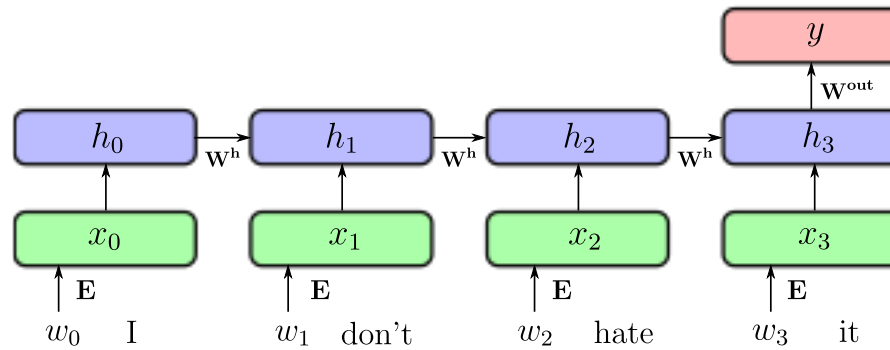
# Backpropagation through time

Similar as standard backpropagation on unrolled network



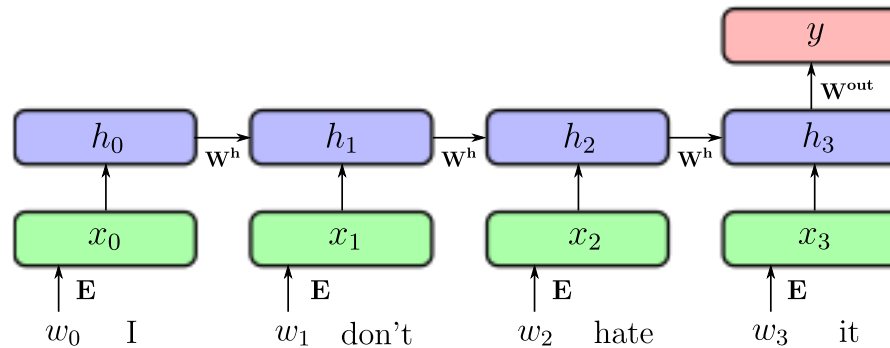
- Similar as training **very deep networks** with tied parameters
- Example between  $x_0$  and  $y_2$ :  $W^h$  is used twice
- Usually truncate the backprop after  $T$  timesteps
- Difficulties to train long-term dependencies

## Other uses: Sentiment Analysis




- Output is sentiment (1 for positive, 0 for negative)

## Other uses: Sentiment Analysis




- Output is sentiment (1 for positive, 0 for negative)
- Very dependent on words order

## Other uses: Sentiment Analysis



- Output is sentiment (1 for positive, 0 for negative)
- Very dependent on words order
- Very flexible network architectures


## Other uses: Sentiment analysis



- Output is sentiment (1 for positive, 0 for negative)
- Very dependent on words order
- Very flexible network architectures




# LSTM




Hochreiter, Sepp, and Jürgen Schmidhuber. "Long short-term memory." Neural computation 1997

# LSTM




Hochreiter, Sepp, and Jürgen Schmidhuber. "Long short-term memory." Neural computation 1997

# LSTM



Hochreiter, Sepp, and Jürgen Schmidhuber. "Long short-term memory." Neural computation 1997


# LSTM



- 4 times more parameters than RNN

Hochreiter, Sepp, and Jürgen Schmidhuber. "Long short-term memory." Neural computation 1997


# LSTM



- 4 times more parameters than RNN
- Mitigates **vanishing gradient** problem through **gating**

Hochreiter, Sepp, and Jürgen Schmidhuber. "Long short-term memory." Neural computation 1997

# LSTM



- 4 times more parameters than RNN
- Mitigates **vanishing gradient** problem through **gating**
- Widely used and SOTA in many sequence learning problems

Hochreiter, Sepp, and Jürgen Schmidhuber. "Long short-term memory." Neural computation 1997

# Vanishing / Exploding Gradients

Passing through  $t$  time-steps, the resulting gradient is the product of many gradients and activations.

# Vanishing / Exploding Gradients

Passing through  $t$  time-steps, the resulting gradient is the product of many gradients and activations.

- Gradient messages close to 0 can shrink to 0
- Gradient messages larger than 1 can explode



# Vanishing / Exploding Gradients

Passing through  $t$  time-steps, the resulting gradient is the product of many gradients and activations.

- Gradient messages close to 0 can shrink to 0
- Gradient messages larger than 1 can explode
- LSTM mitigates that in RNNs
- Additive path between  $c_t$  and  $c_{t-1}$

# Vanishing / Exploding Gradients

Passing through  $t$  time-steps, the resulting gradient is the product of many gradients and activations.

- Gradient messages close to 0 can shrink to 0
- Gradient messages larger than 1 can explode
- LSTM mitigates that in RNNs
- Additive path between  $c_t$  and  $c_{t-1}$
- Gradient clipping prevents gradient explosion
- Well chosen activation function is critical (tanh)

# Vanishing / Exploding Gradients

Passing through  $t$  time-steps, the resulting gradient is the product of many gradients and activations.

- Gradient messages close to 0 can shrink to 0
- Gradient messages larger than 1 can explode
- LSTM mitigates that in RNNs
- Additive path between  $c_t$  and  $c_{t-1}$
- Gradient clipping prevents gradient explosion
- Well chosen activation function is critical (tanh)

Skip connections in ResNet also alleviate a similar optimization problem.

Next: Lab 6!