Deep Learning

GPT and Attention

What is Natural Language Processing (NLP)?

- Subfield of AI that focuses on reading, deciphering and producing human language
- Combines computational linguistics (e.g. rule-based modelling of language) with statistical, ML, and deep learning approaches
- Through NLP, machines can understand, analyze and generate language in ways that are meaningful and contextually appropriate
- While LLMs have driven an explosion in interest, there are many older technologies which paved the way

Common Applications

- **Text Classification:** e.g. spam detection, news classification
- **Question Answering:** processing input and finding relevant information
- **Machine Translation**
- **Sentiment Analysis**
- **Code Generation**

Challenges in NLP

- **Ambiguity:** Words can mean different things depending on context
- **Nuances:** Languages are full of idioms, slang, cultural references, sarcasm…
- **Syntax vs Semantics:** A grammatically correct sentence might not make sense, or a grammatically incorrect one might be easy to understand

- I saw a man on the hill with the telescope
- That’s a cool cat
- Colourless green ideas sleep furiously
- Me went store

Challenges in NLP

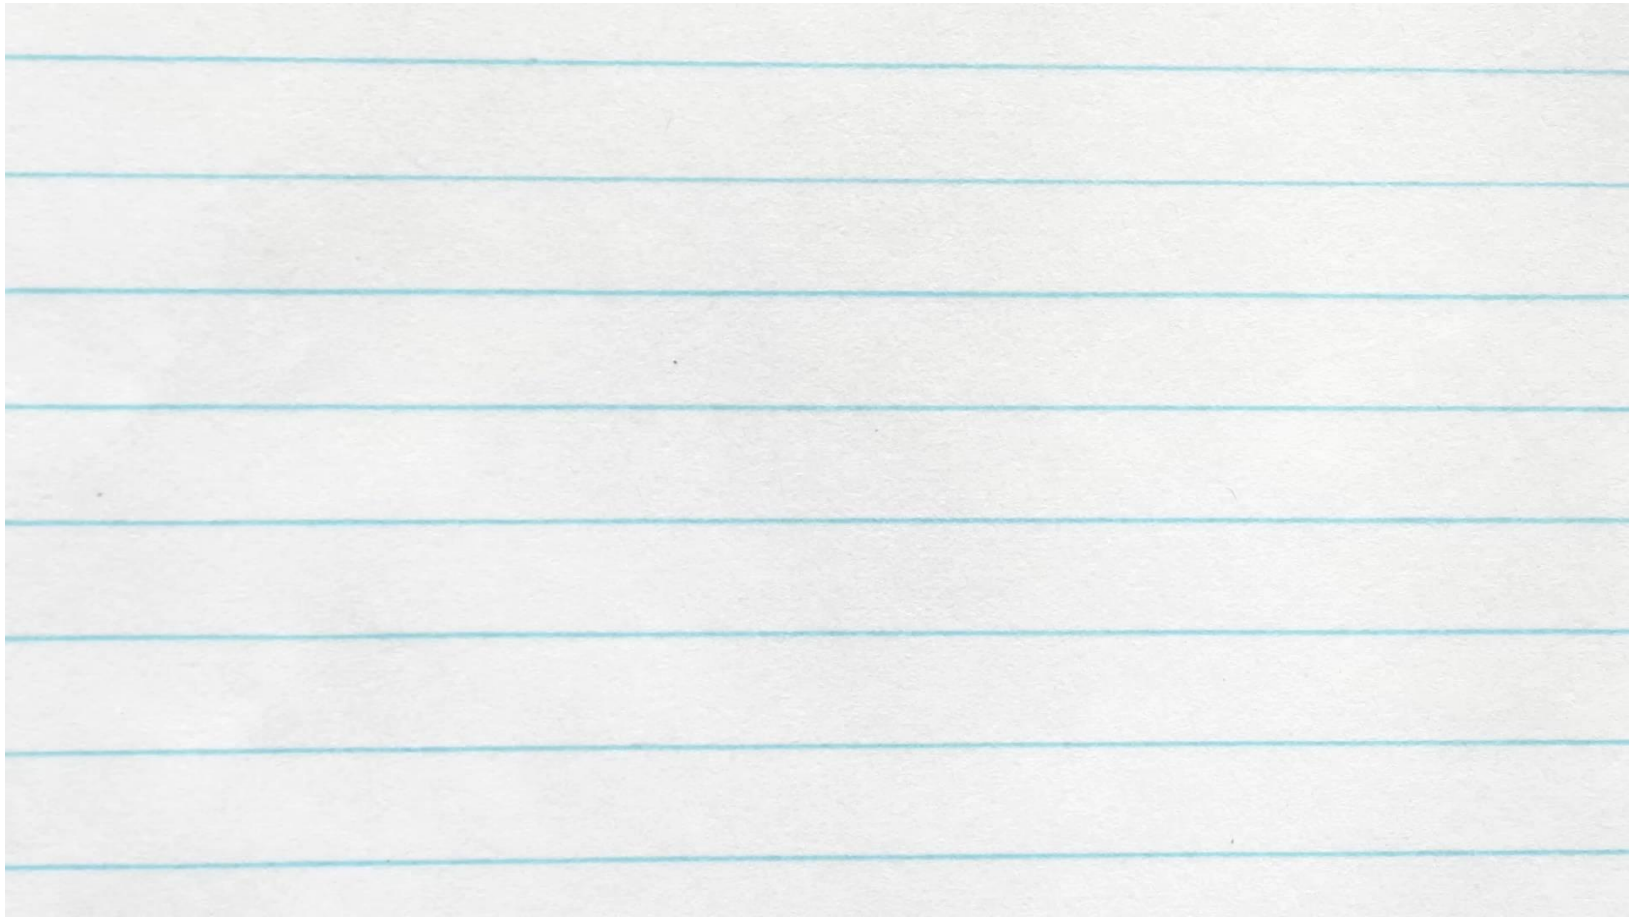

Quick history of NLP

- **1950s:**
  - Alan Turing publishes “Computing Machinery and Intelligence”, in which he proposes the Turing test
  - Noam Chomsky publishes “Syntactic Structures”, an attempt to construct a formal theory of linguistic structure
- **1960s:**
  - Georgetown University develops a machine translation system, which automatically translates 60 Russian sentences into English using an extremely complex flowchart and a limited vocabulary
  - The authors claim machine translation could be a solved problem in five years

Quick history of NLP

- **1980s – 1990s:**
  - Transition from rule-based to statistical approaches
  - Idea of using existing text to train a model begins to appear (e.g. bilingual documents from the Canadian parliament)
- **2000s – 2010s:**
  - Deep Learning takes over
  - Tools such as convolutional networks, and later RNNs, transform the field
- **2020s: The era of the Large Language Model**

Foundations of LLMs

- Fundamental goal of language modelling: next word prediction
- $P(cat \mid the\ dog\ and\ the)$
- To generate, pick the word with highest likelihood
- Early models could handle one, two words of context
- Locally coherent, but longer texts quickly lose meaning
- More context requires more complexity!

How does an LM “understand” word meaning?

- In order to predict the likelihood of a word, we must have some sense of its meaning
- Some words have similar meanings, and can easily fit in the same place
- In the same way CNNs convert an image into a set of feature maps, we can convert a word into a set of abstract linguistic features
- Word2Vec: 300 features
- GPT-3: 12,888

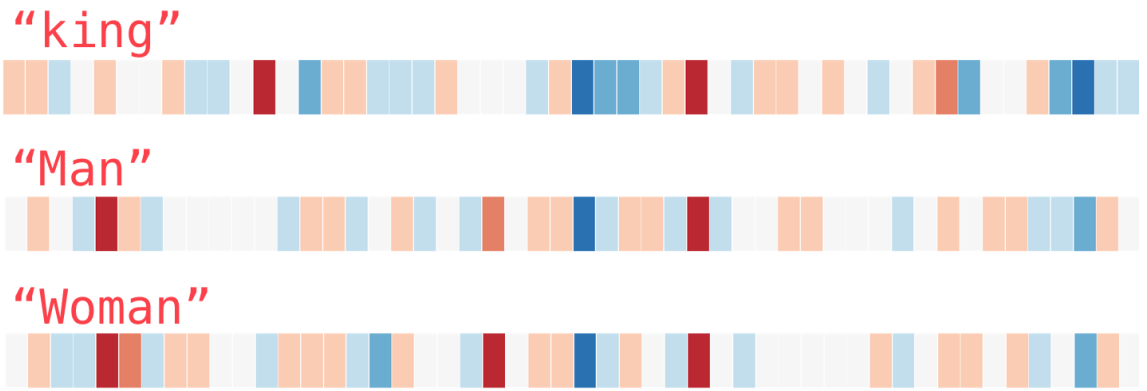

How does an LM “understand” word meaning?

- Key concept of word embeddings: similar words should have similar vectors
- How do we accomplish this?

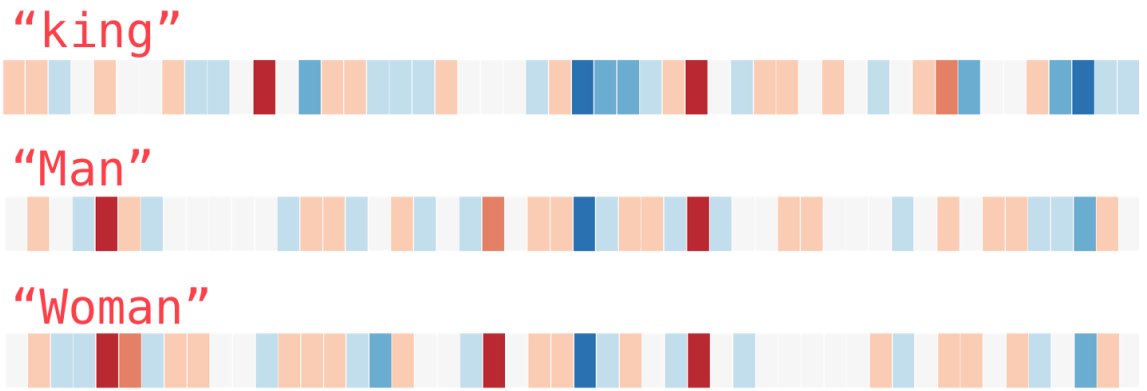

Building Word Embeddings

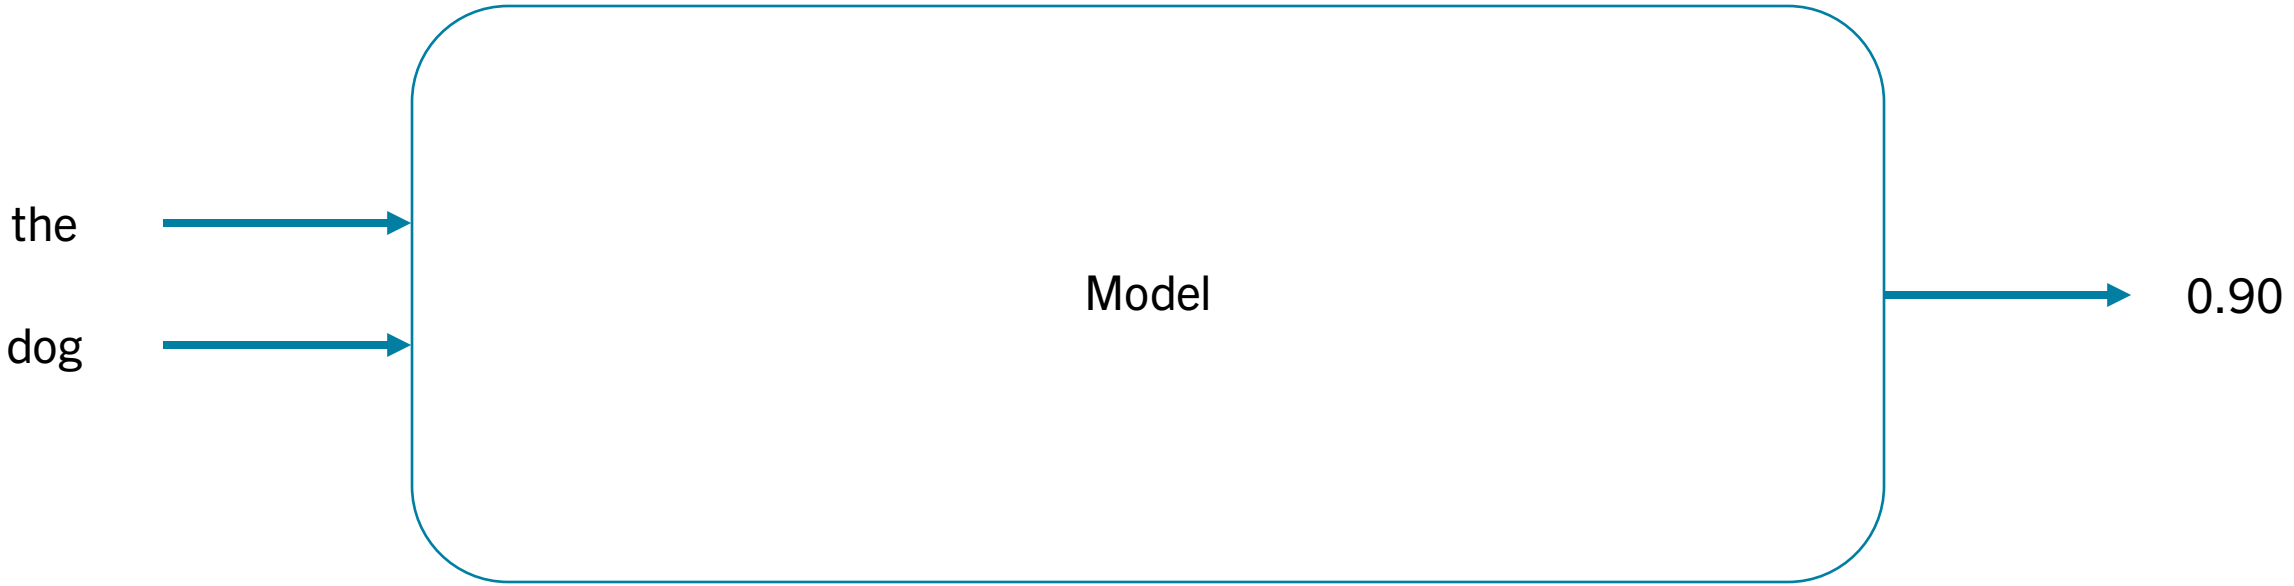

Building Word Embeddings

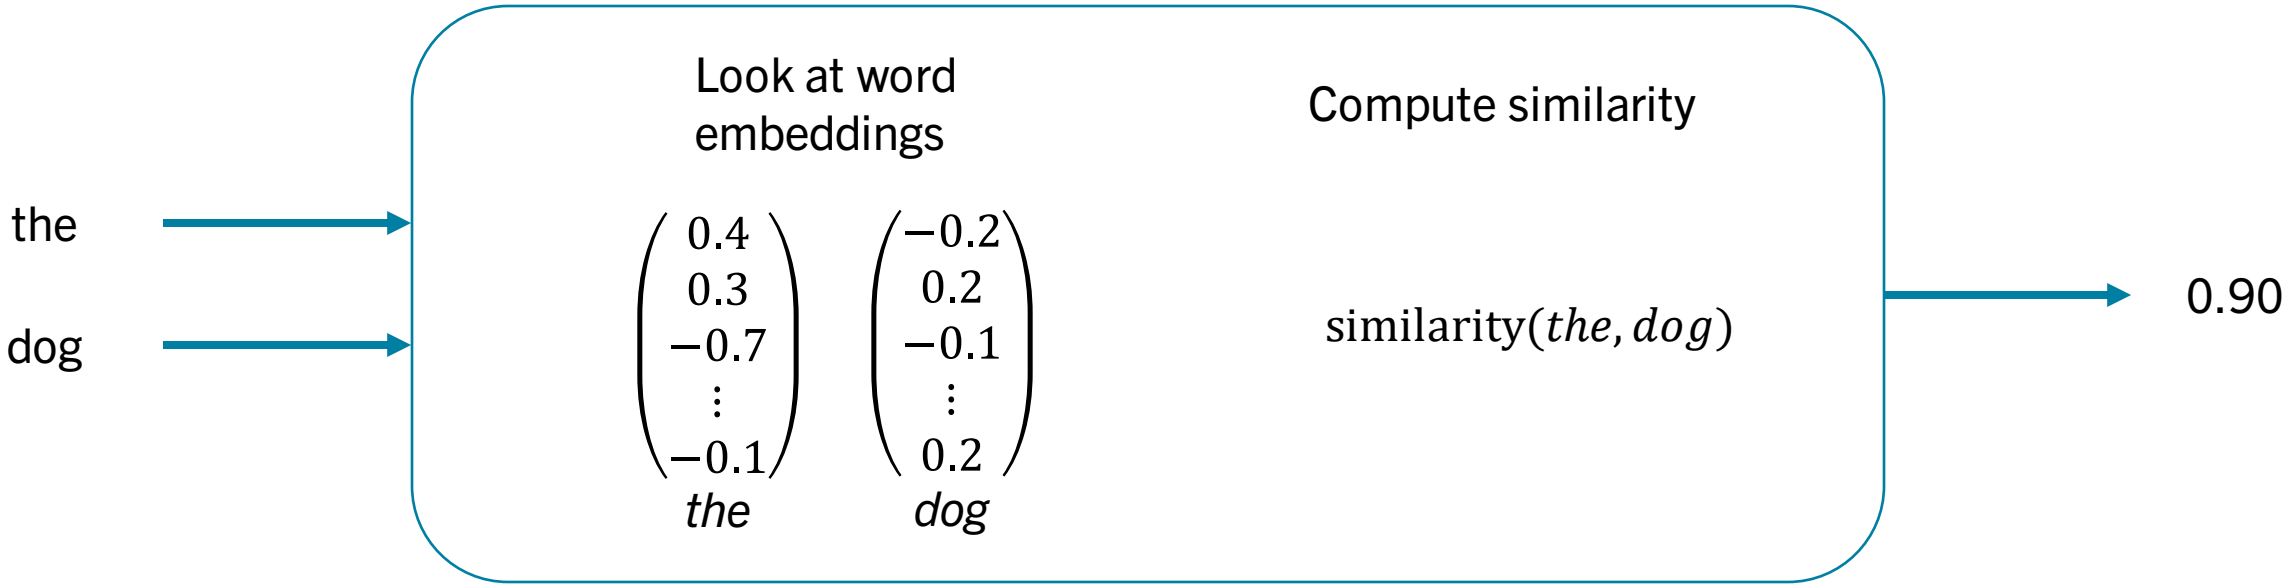

Building Word Embeddings

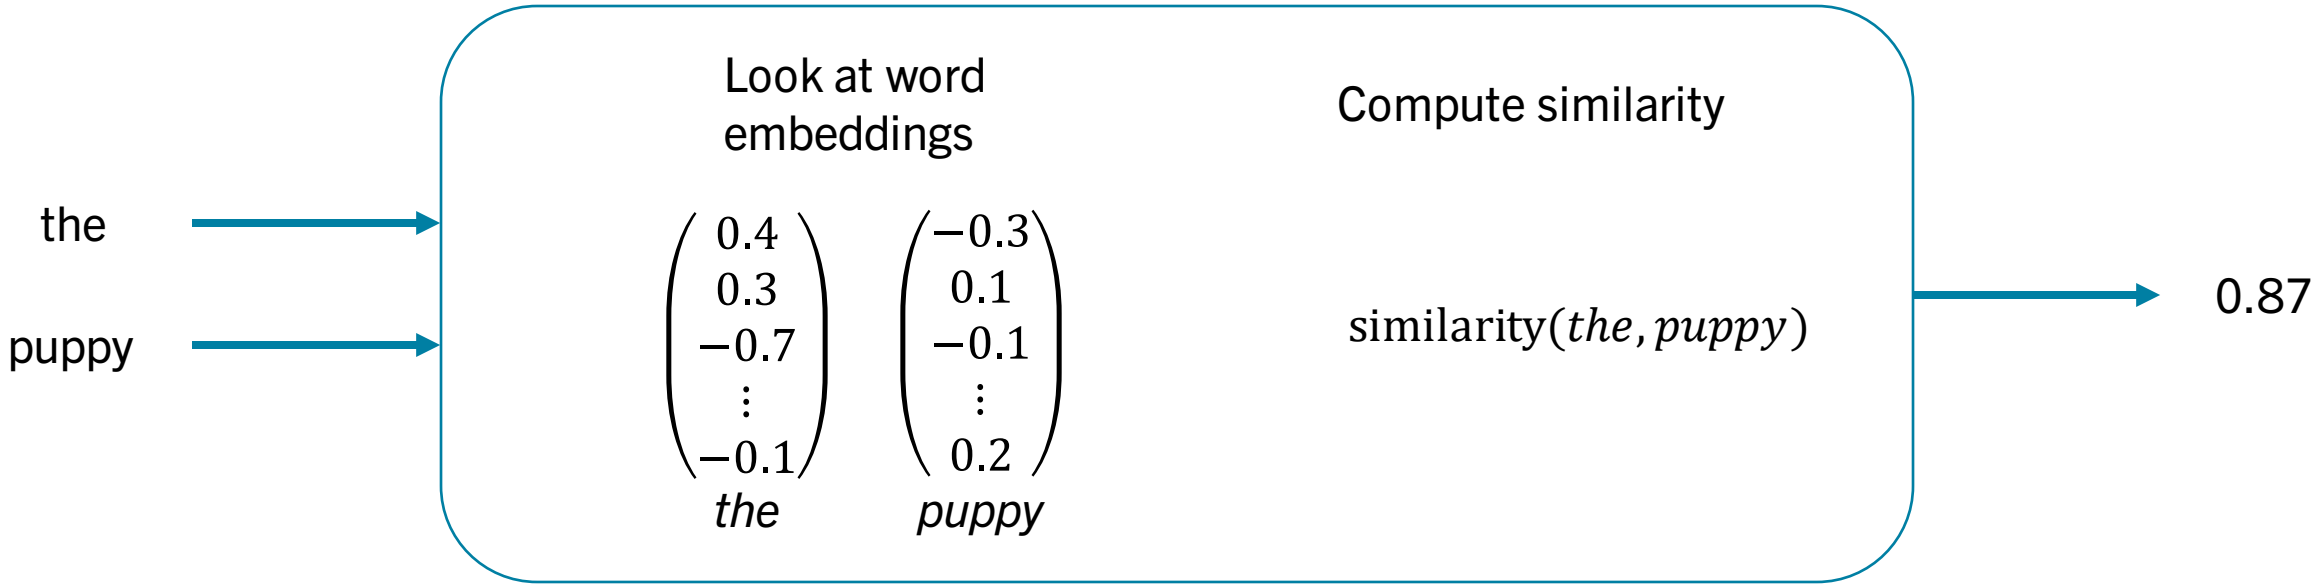

Building Word Embeddings

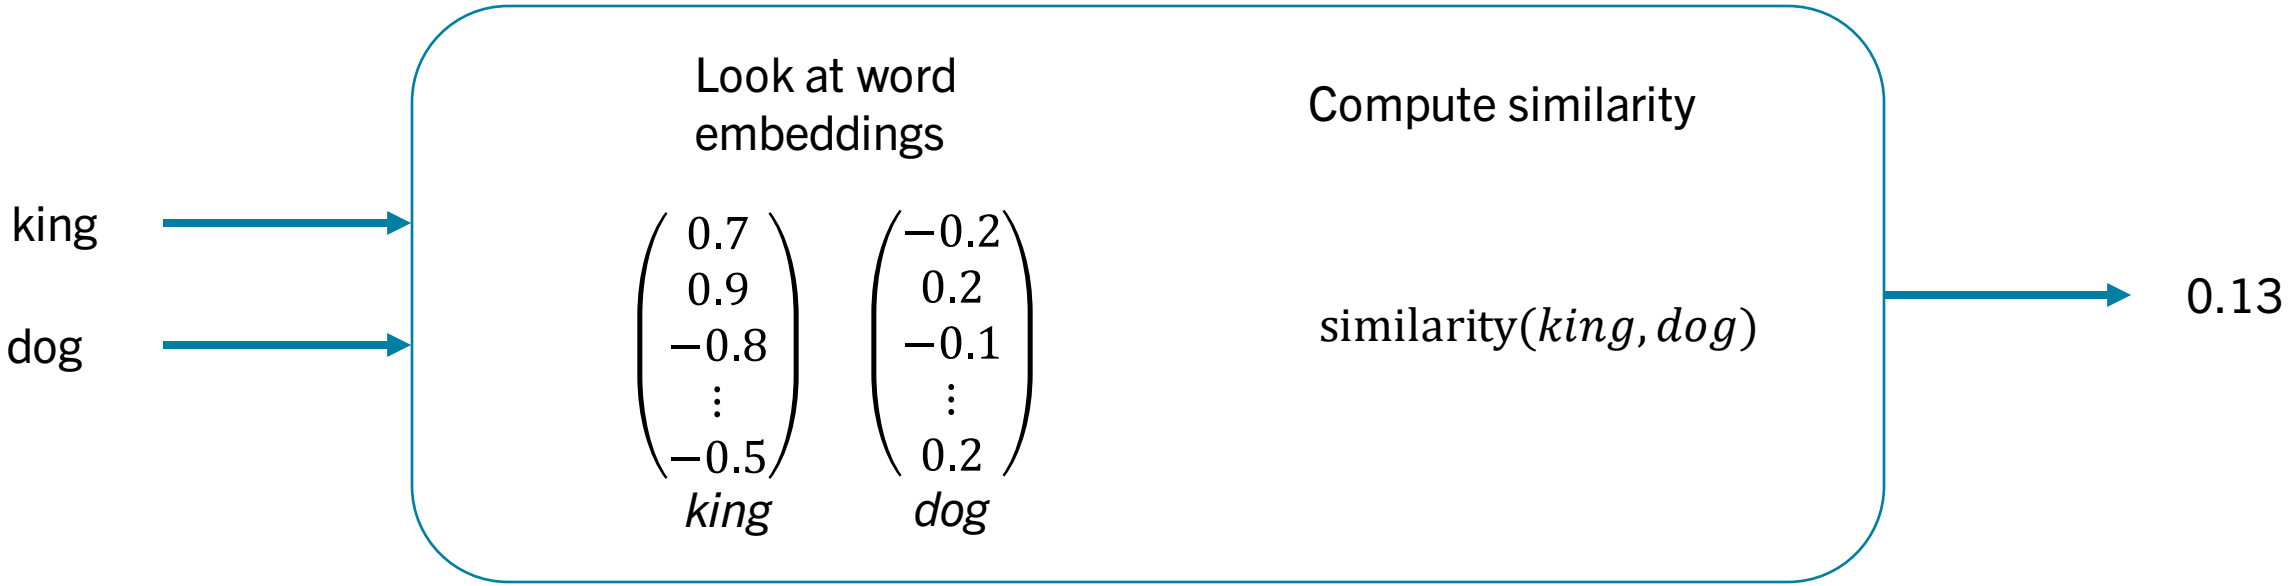

Results

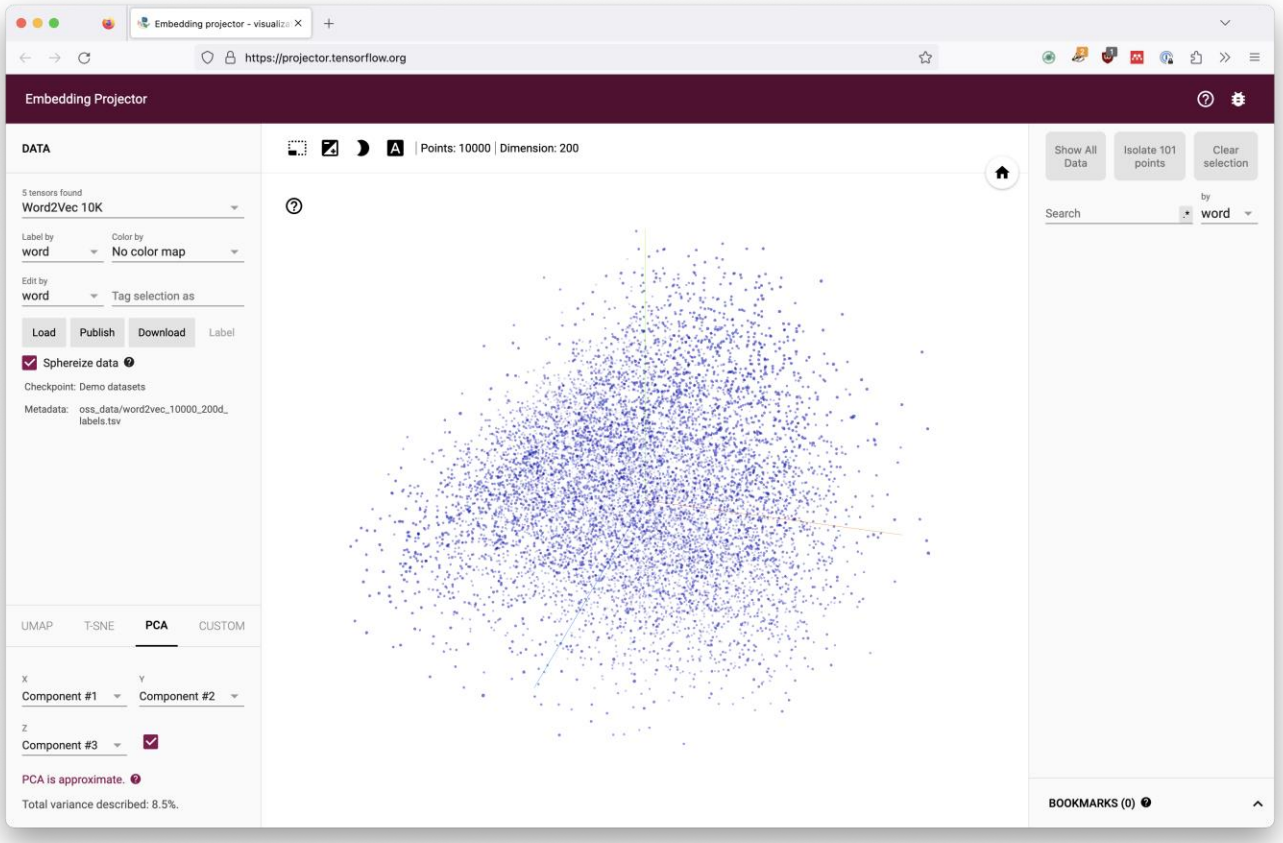
Embedding projector - visualize
https://projector.tensorflow.org
Embedding Projector
DATA
5 tensors found
Word2Vec 10K
Label by
word
Color by
No color map
Edit by
word
Tag selection as
Load
Publish
Download
Label
Sphereize data
Checkpoint: Demo datasets
Metadata: oss_data/word2vec_10000_200d_labels.tsv
UMAP
T-SNE
PCA
CUSTOM
X
Component #1
Y
Component #2
Z
Component #3
PCA is approximate.
Total variance described: 8.5%.
Points: 10000 | Dimension: 200
Show All Data
Isolate 101 points
Clear selection
Search
by
word
BOOKMARKS (0)

Results

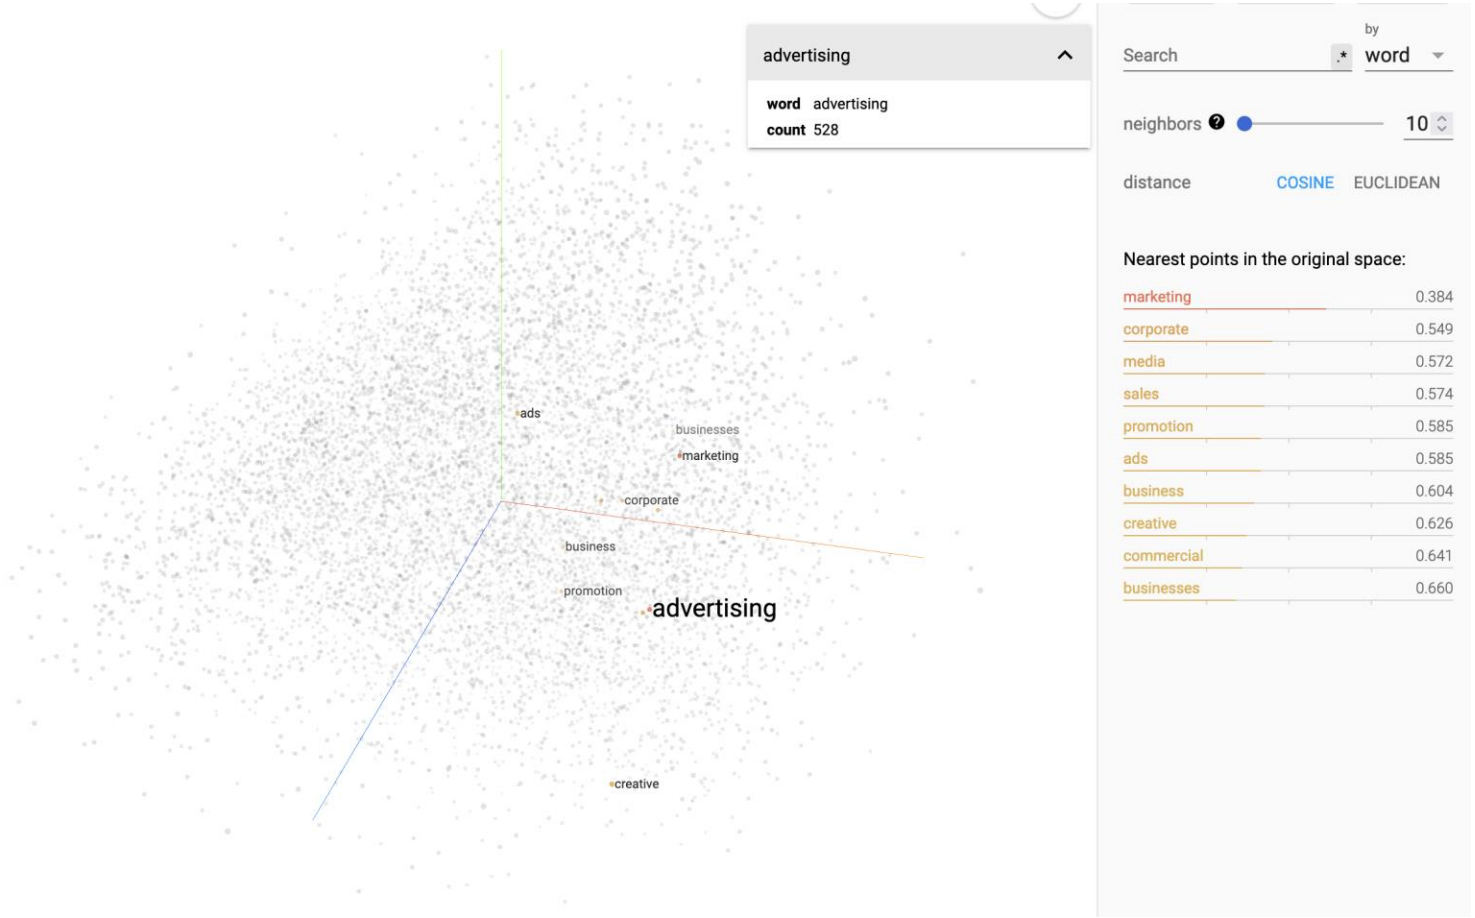
advertising
word advertising
count 528
Search
by
word
neighbors
10
distance
COSINE
EUCLIDEAN
Nearest points in the original space:
marketing 0.384
corporate 0.549
media 0.572
sales 0.574
promotion 0.585
ads 0.585
business 0.604
creative 0.626
commercial 0.641
businesses 0.660
ads
businesses
marketing
corporate
business
promotion
advertising
creative

Building GPT

Building GPT: The Transfomer

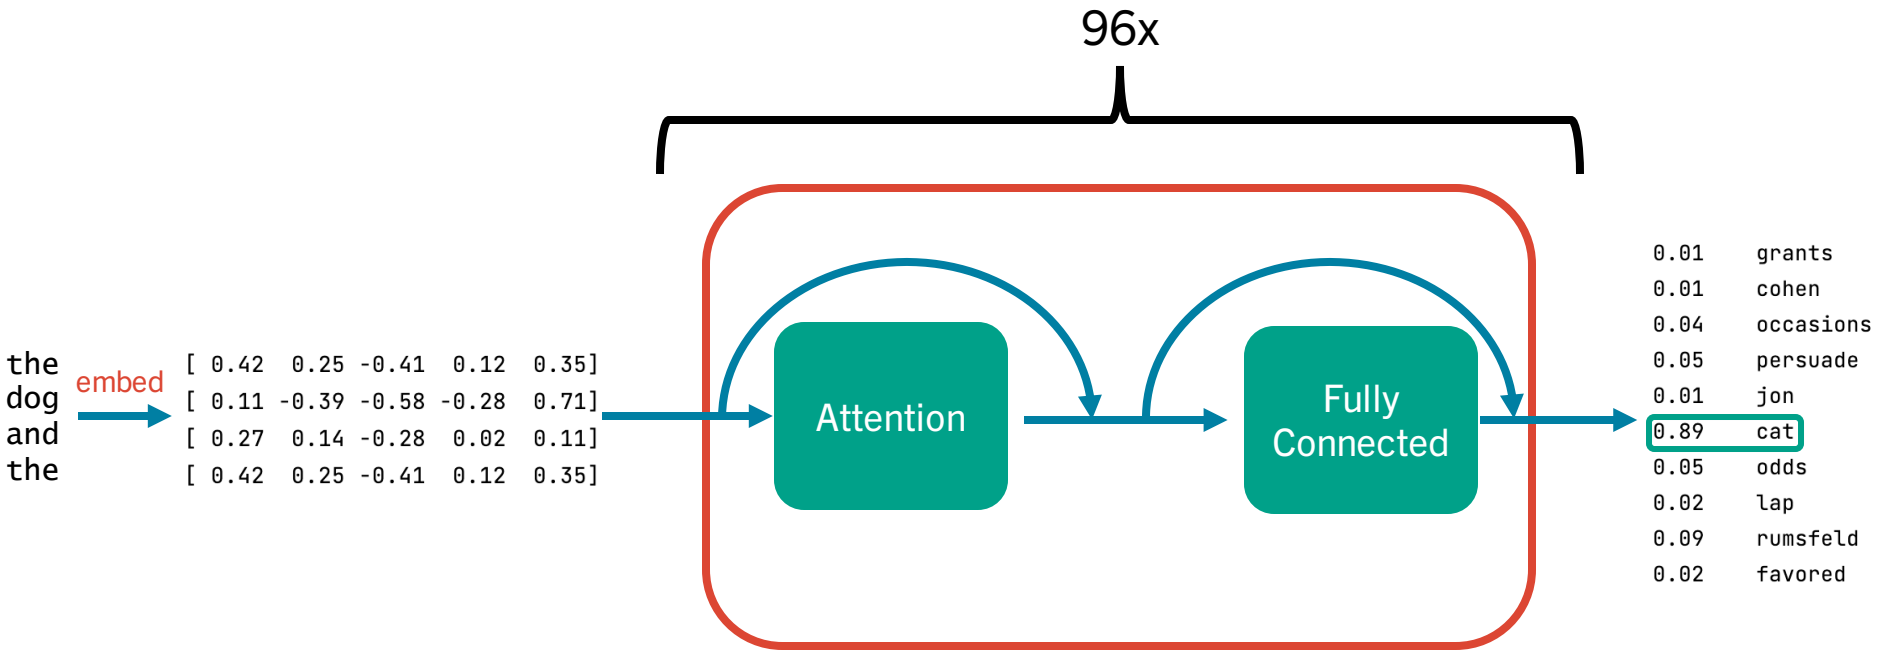

Building GPT: Positional Embedding

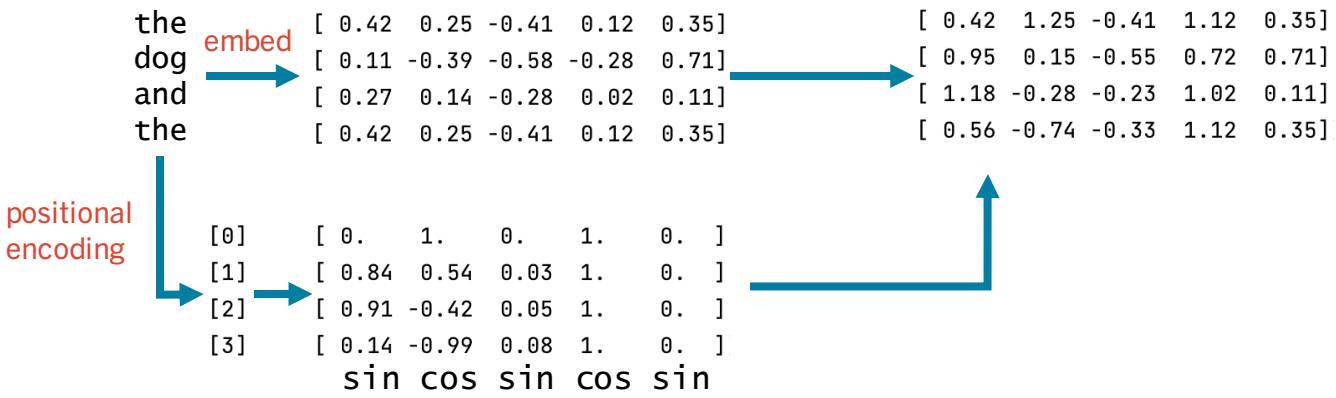

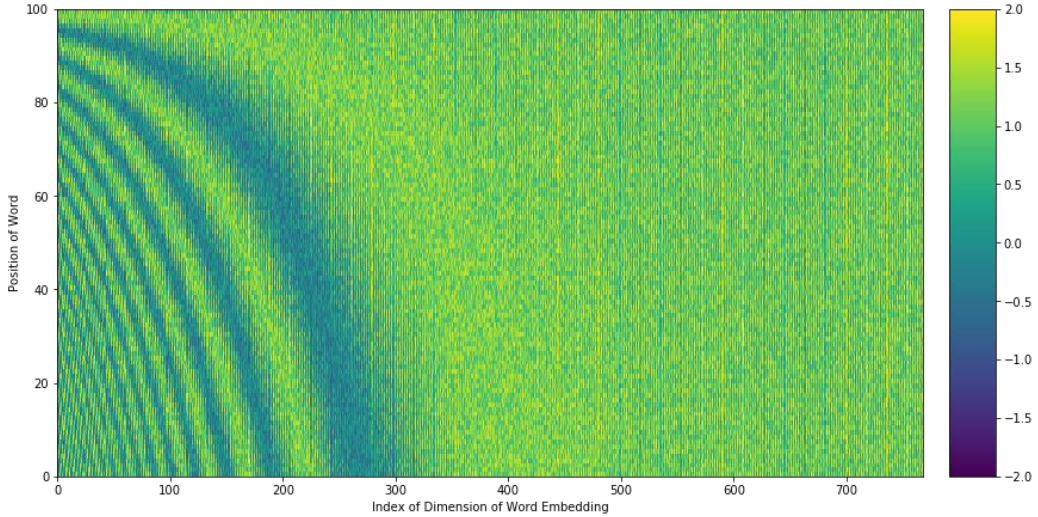

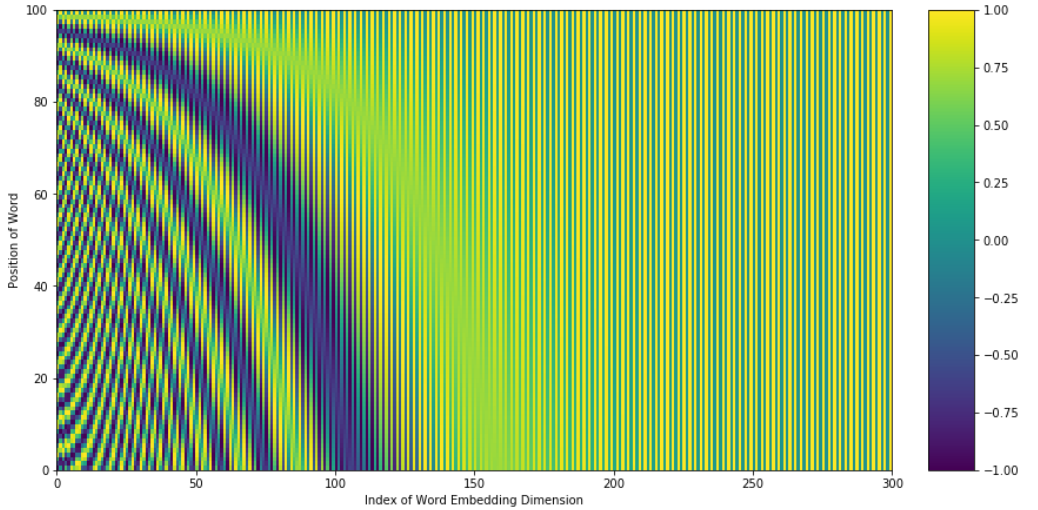

Building GPT

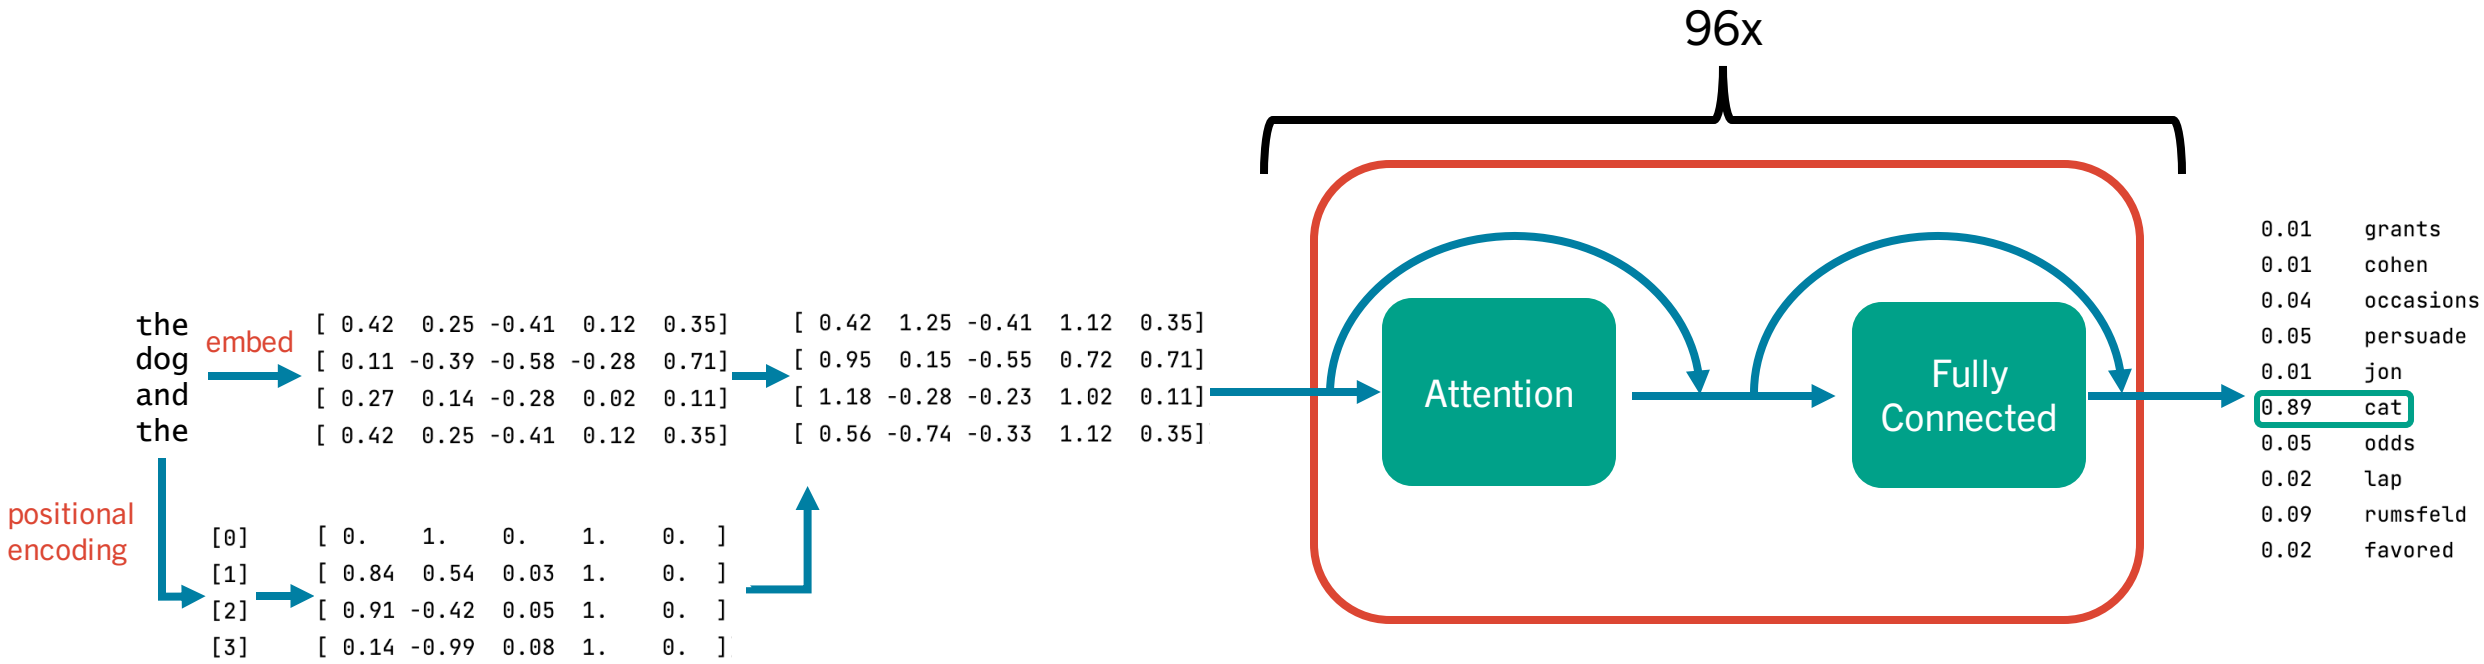

Building GPT

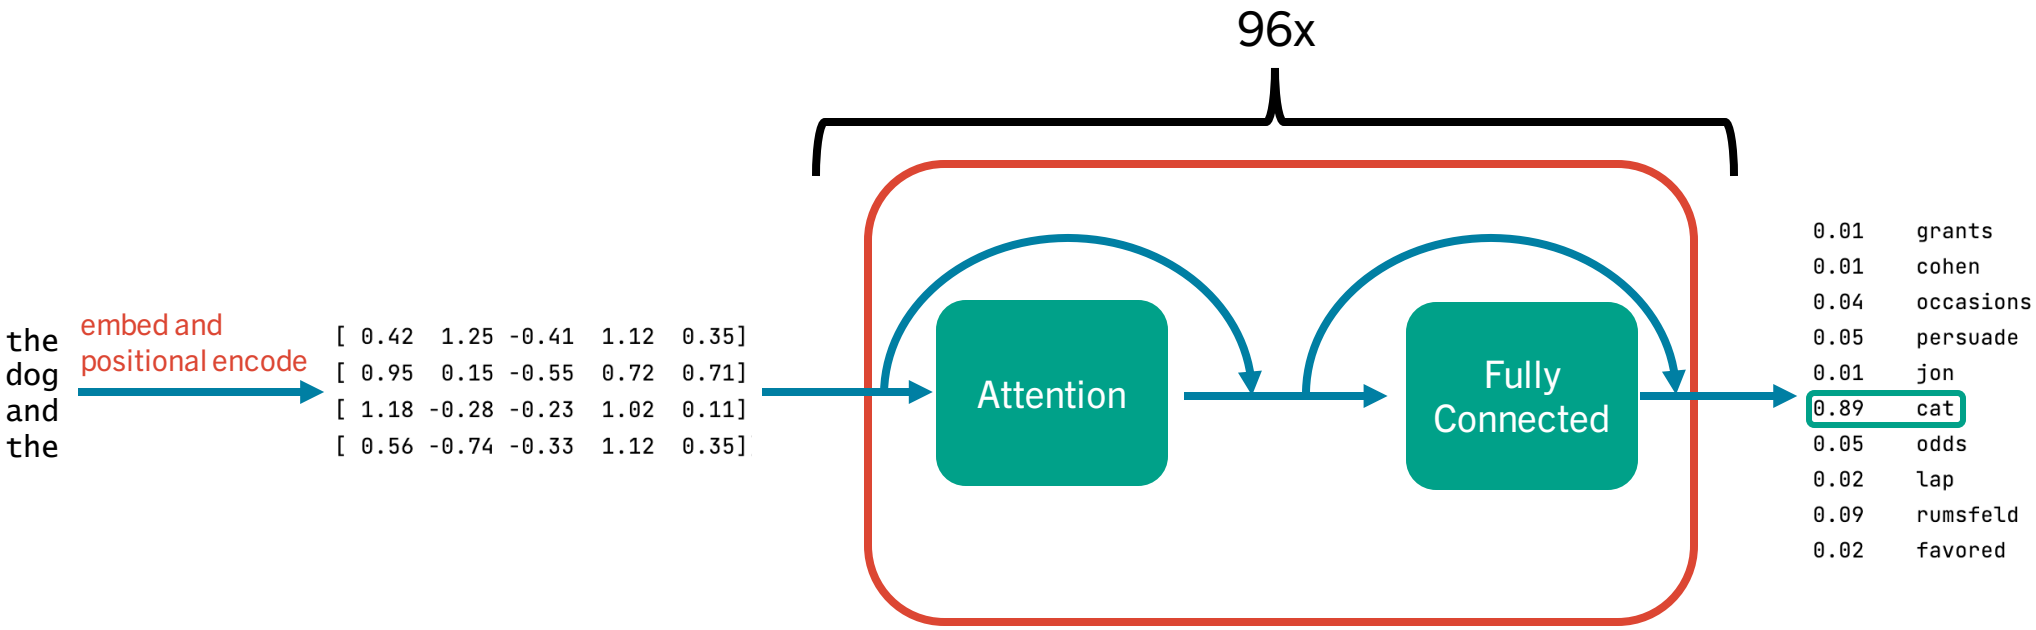

Building GPT: Attention

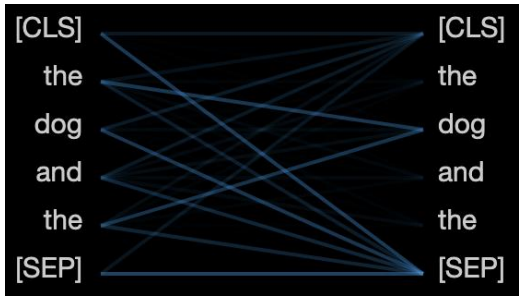

the [0.42 1.25 -0.41 1.12 0.35]
dog [0.95 0.15 -0.55 0.72 0.71]
and [1.18 -0.28 -0.23 1.02 0.11]
the [0.56 -0.74 -0.33 1.12 0.35]

w_Q

w_K

w_V

Query

[0.98 0.18 0.11]
[0.7 0.29 0.72]
[0.42 0.53 0.95]
[0.58 0.06 0.66]

Key

[0.26 0.31 0.22]
[0.81 0.93 0.47]
[0.45 0.49 0.36]
[0.3 0.7 0.79]

Value

[0.8 0.17 0.81]
[0.14 0.01 0.85]
[0.06 0.77 0.27]
[0.37 0.56 0.2]

$softmax(\frac{QK^T}{\sqrt{D_K}})$

	the	dog	and	the
the	[0.21	0.31	0.24	0.23]
dog	[0.2	0.3	0.23	0.27]
and	[0.19	0.3	0.22	0.29]
the	[0.21	0.28	0.24	0.27]

[0.31 0.36 0.55]
[0.31 0.37 0.53]
[0.31 0.37 0.52]
[0.32 0.37 0.53]

Building GPT: Attention

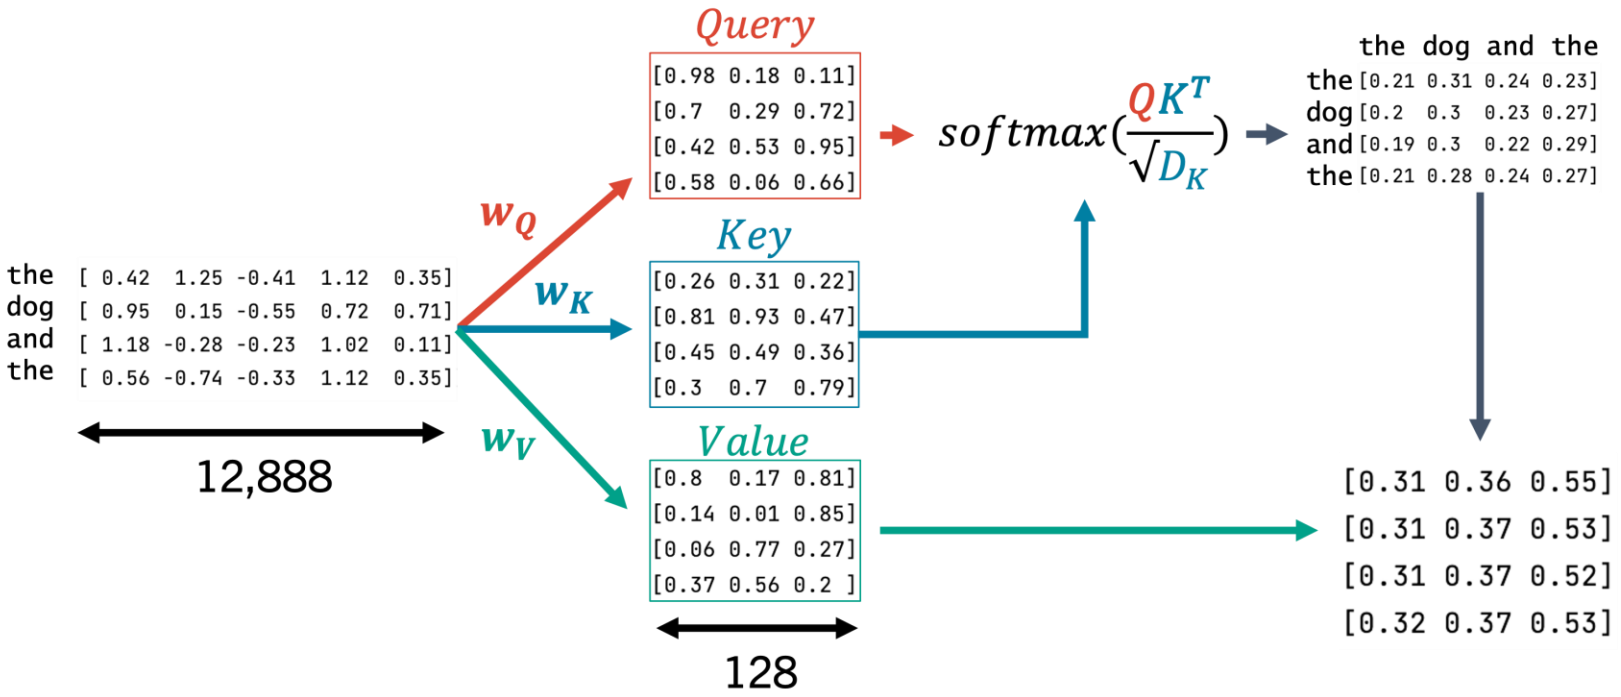

Building GPT: Attention

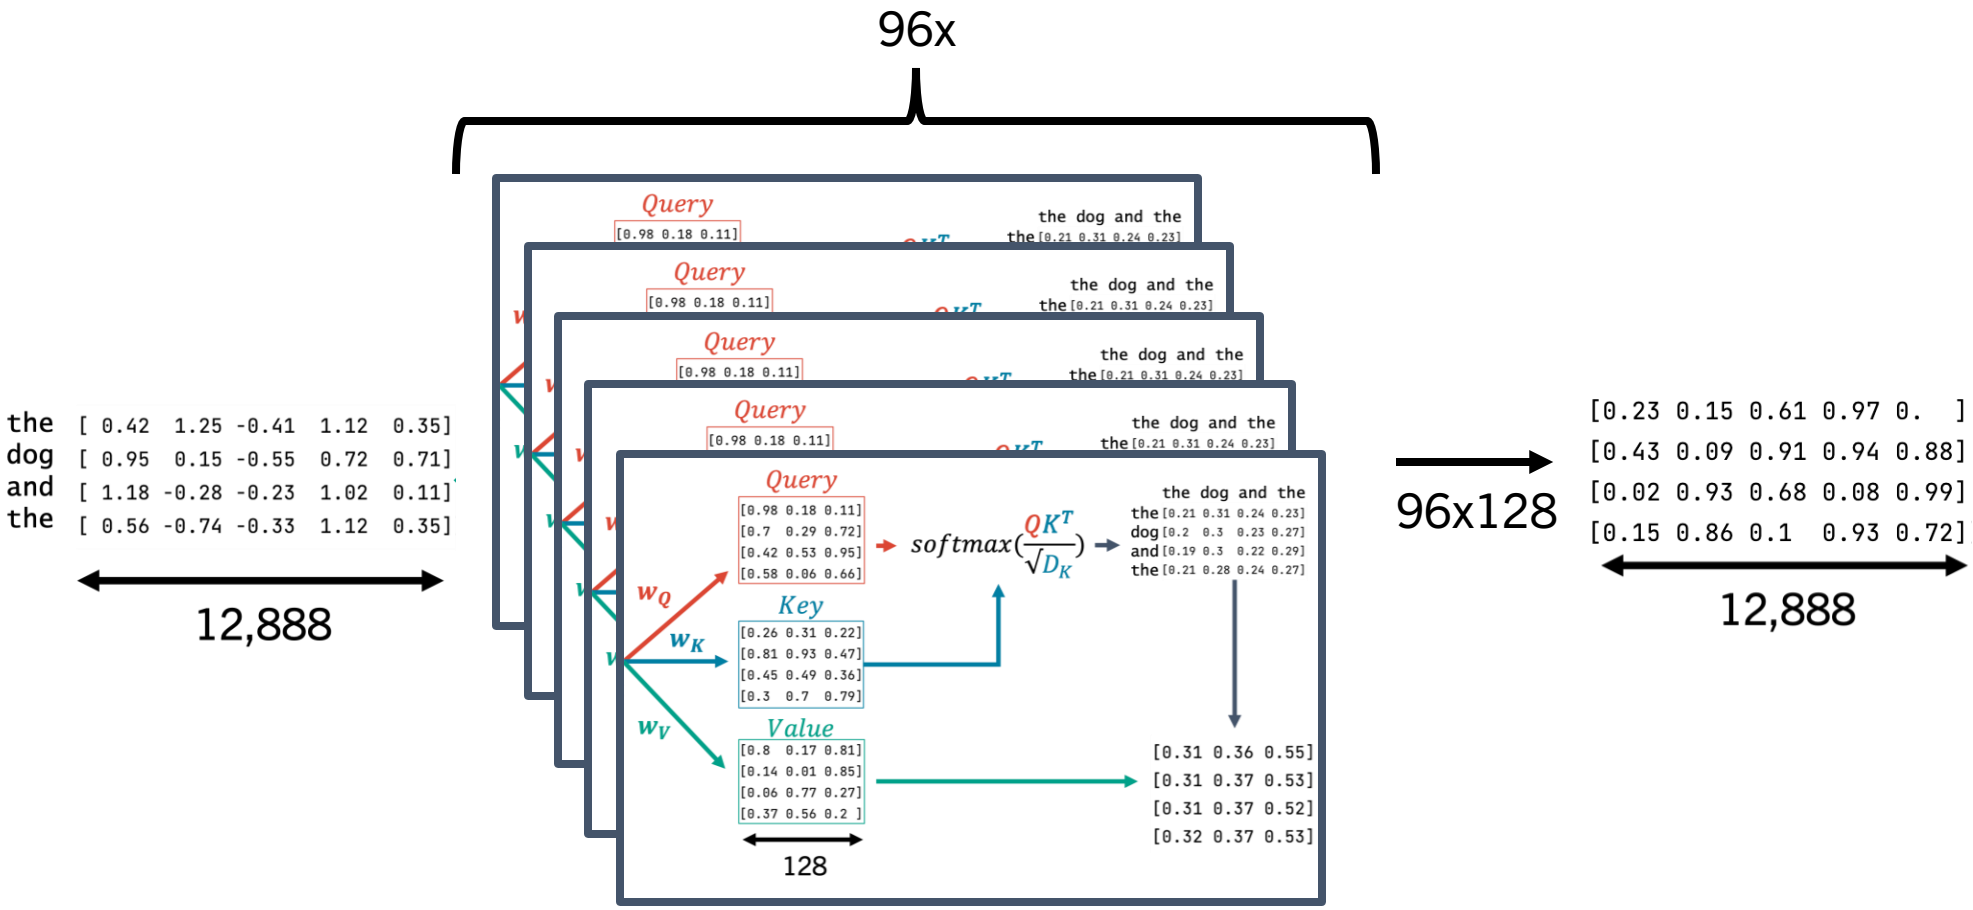

Building GPT

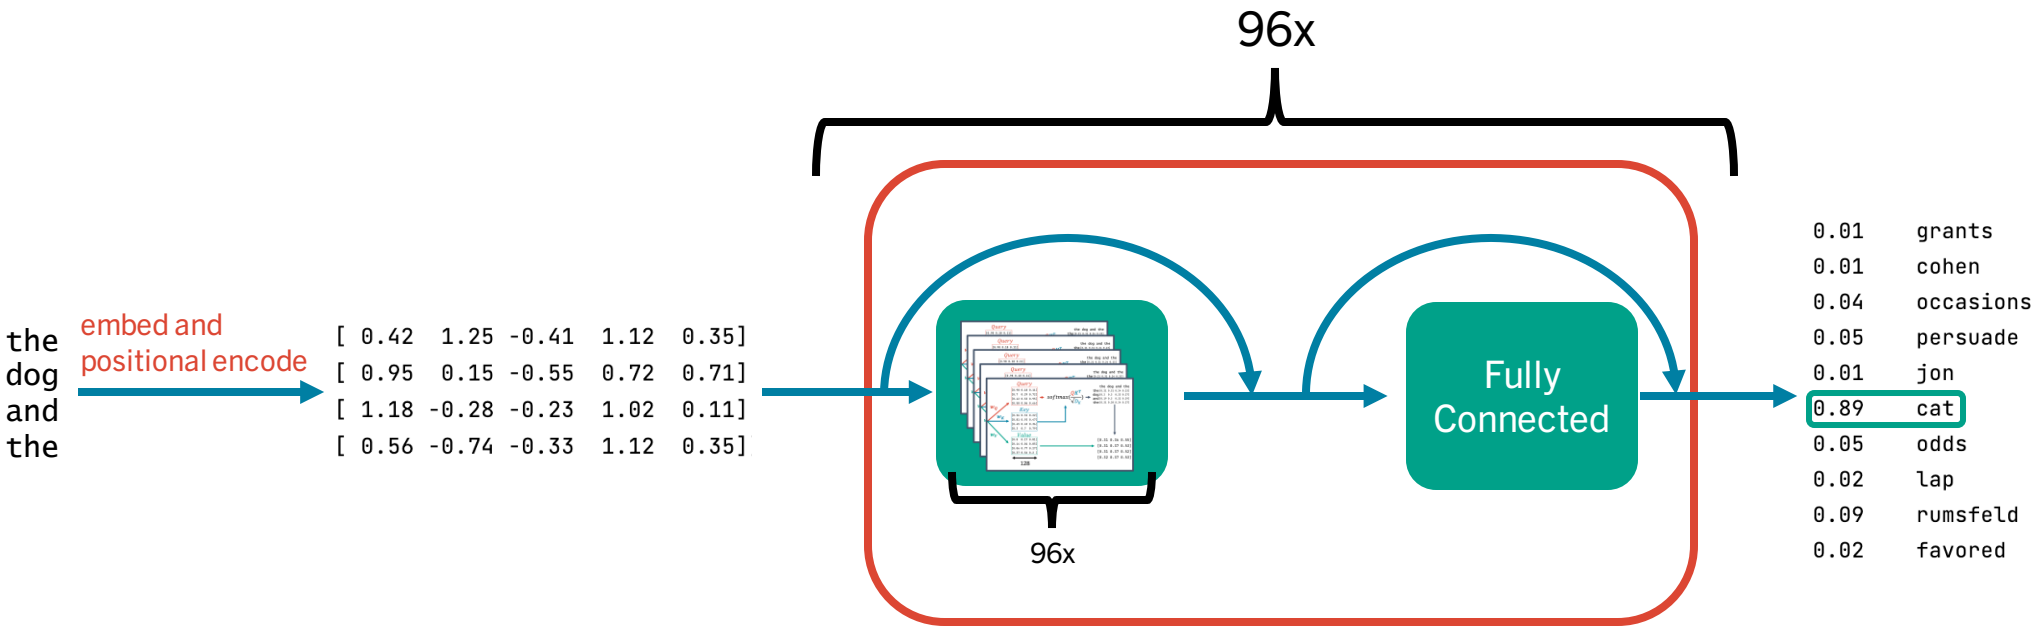

Building GPT: Top-P

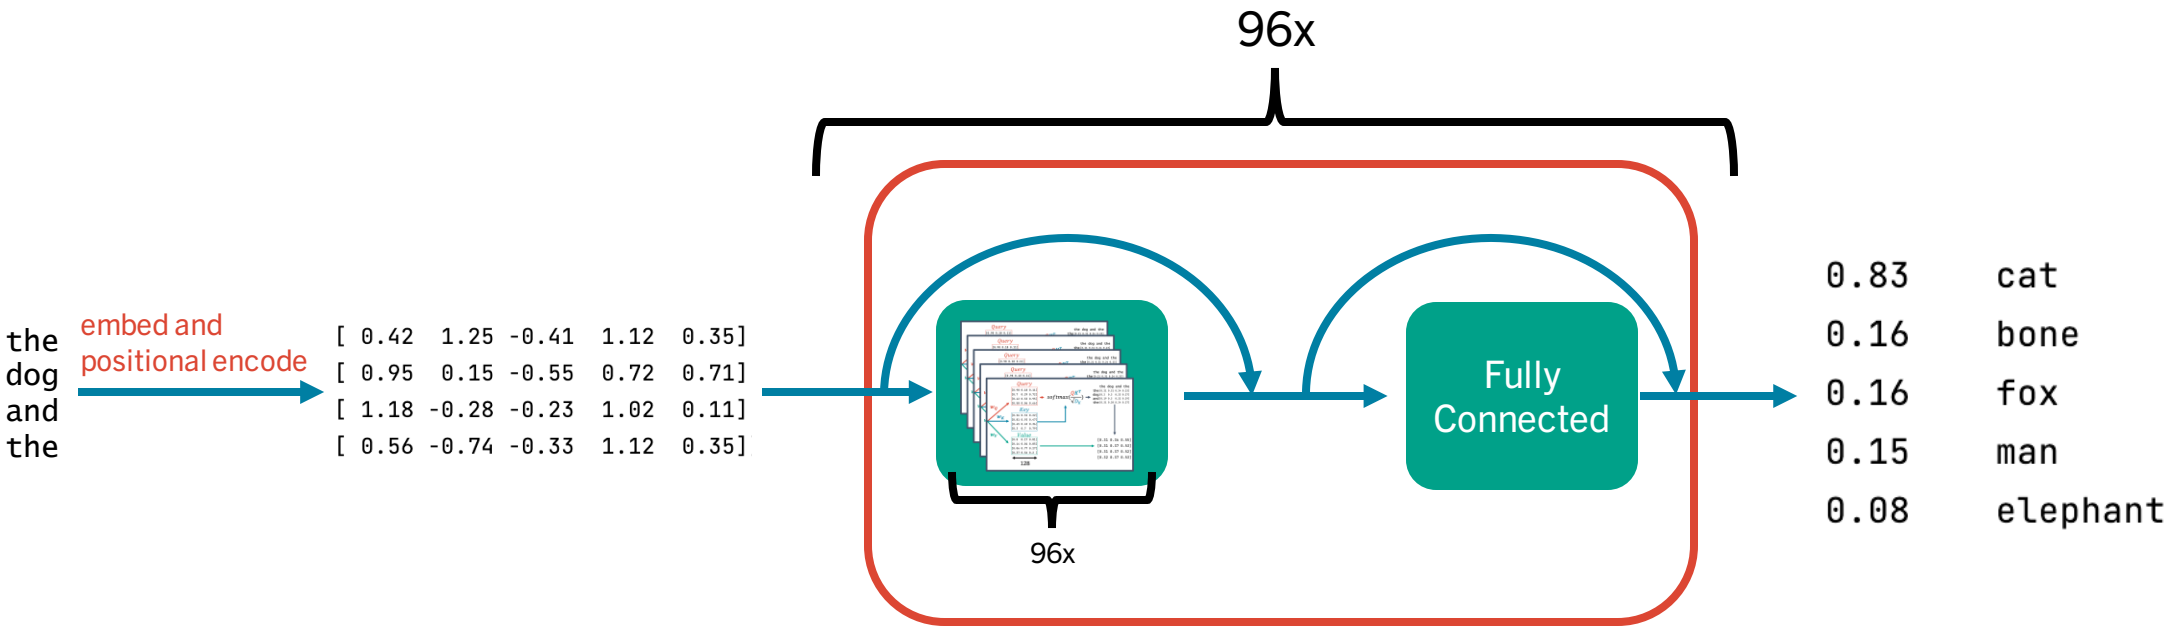

Building GPT: Top-P

Top 10 documentaries about artificial intelligence:

1. AlphaGo (2017)

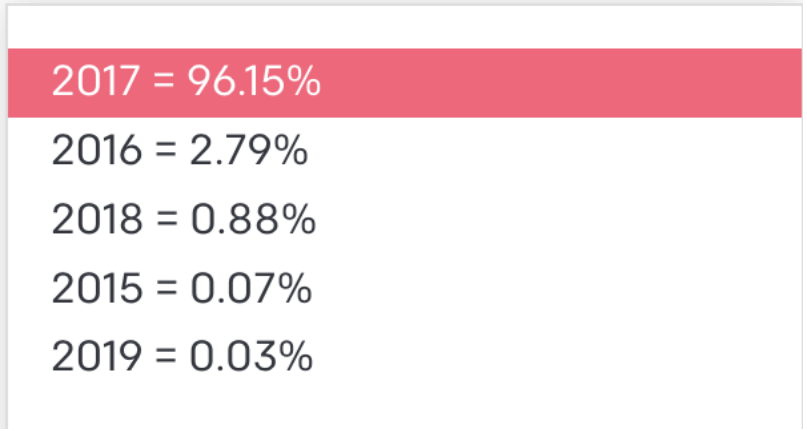

Building GPT

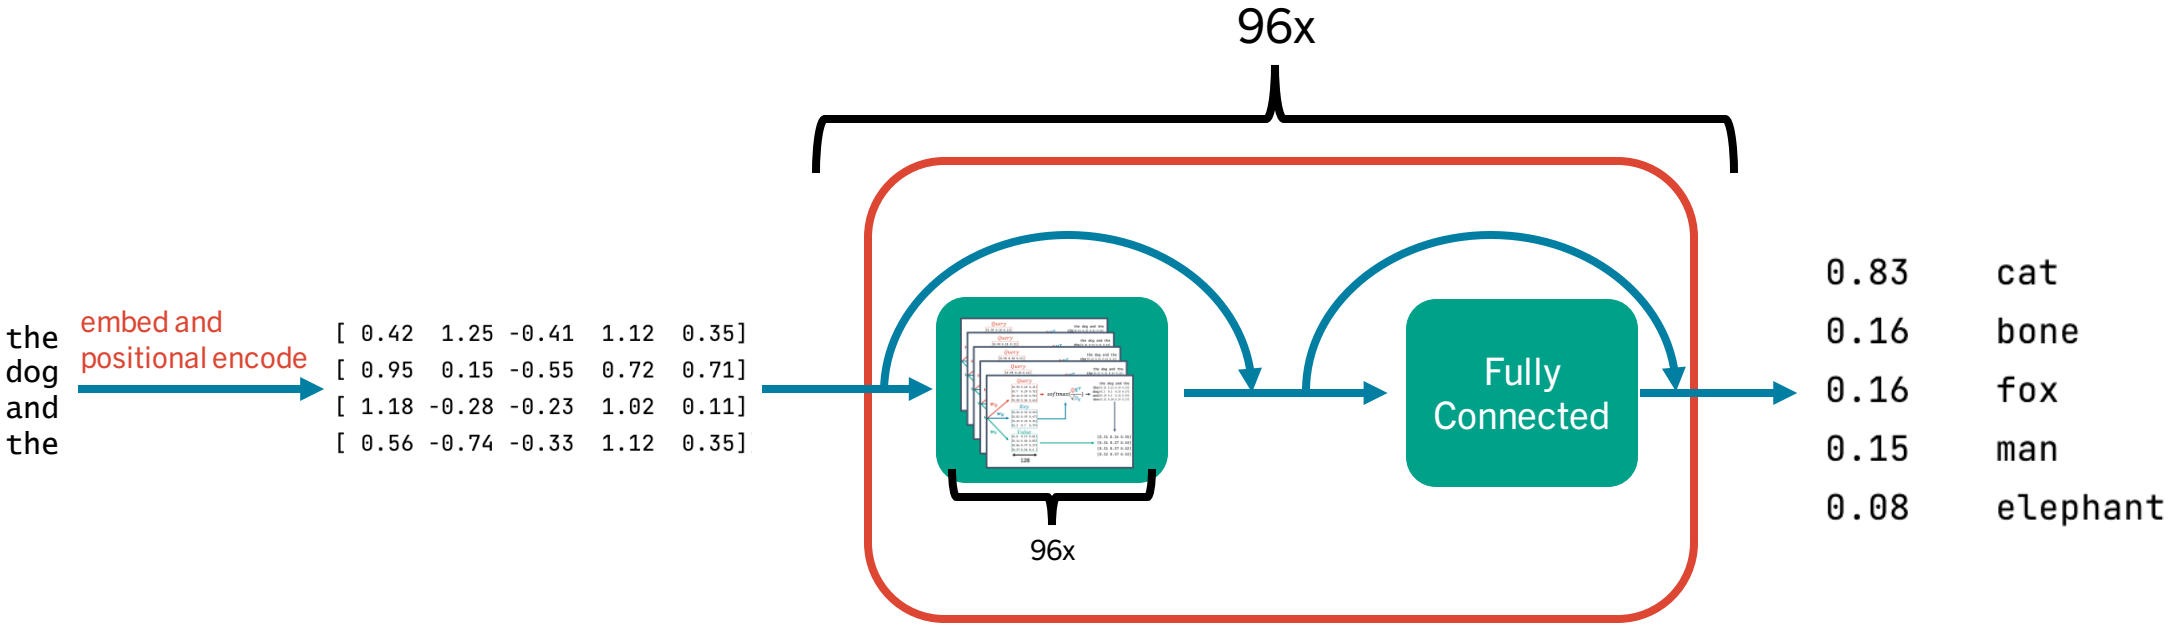

Scale of GPT

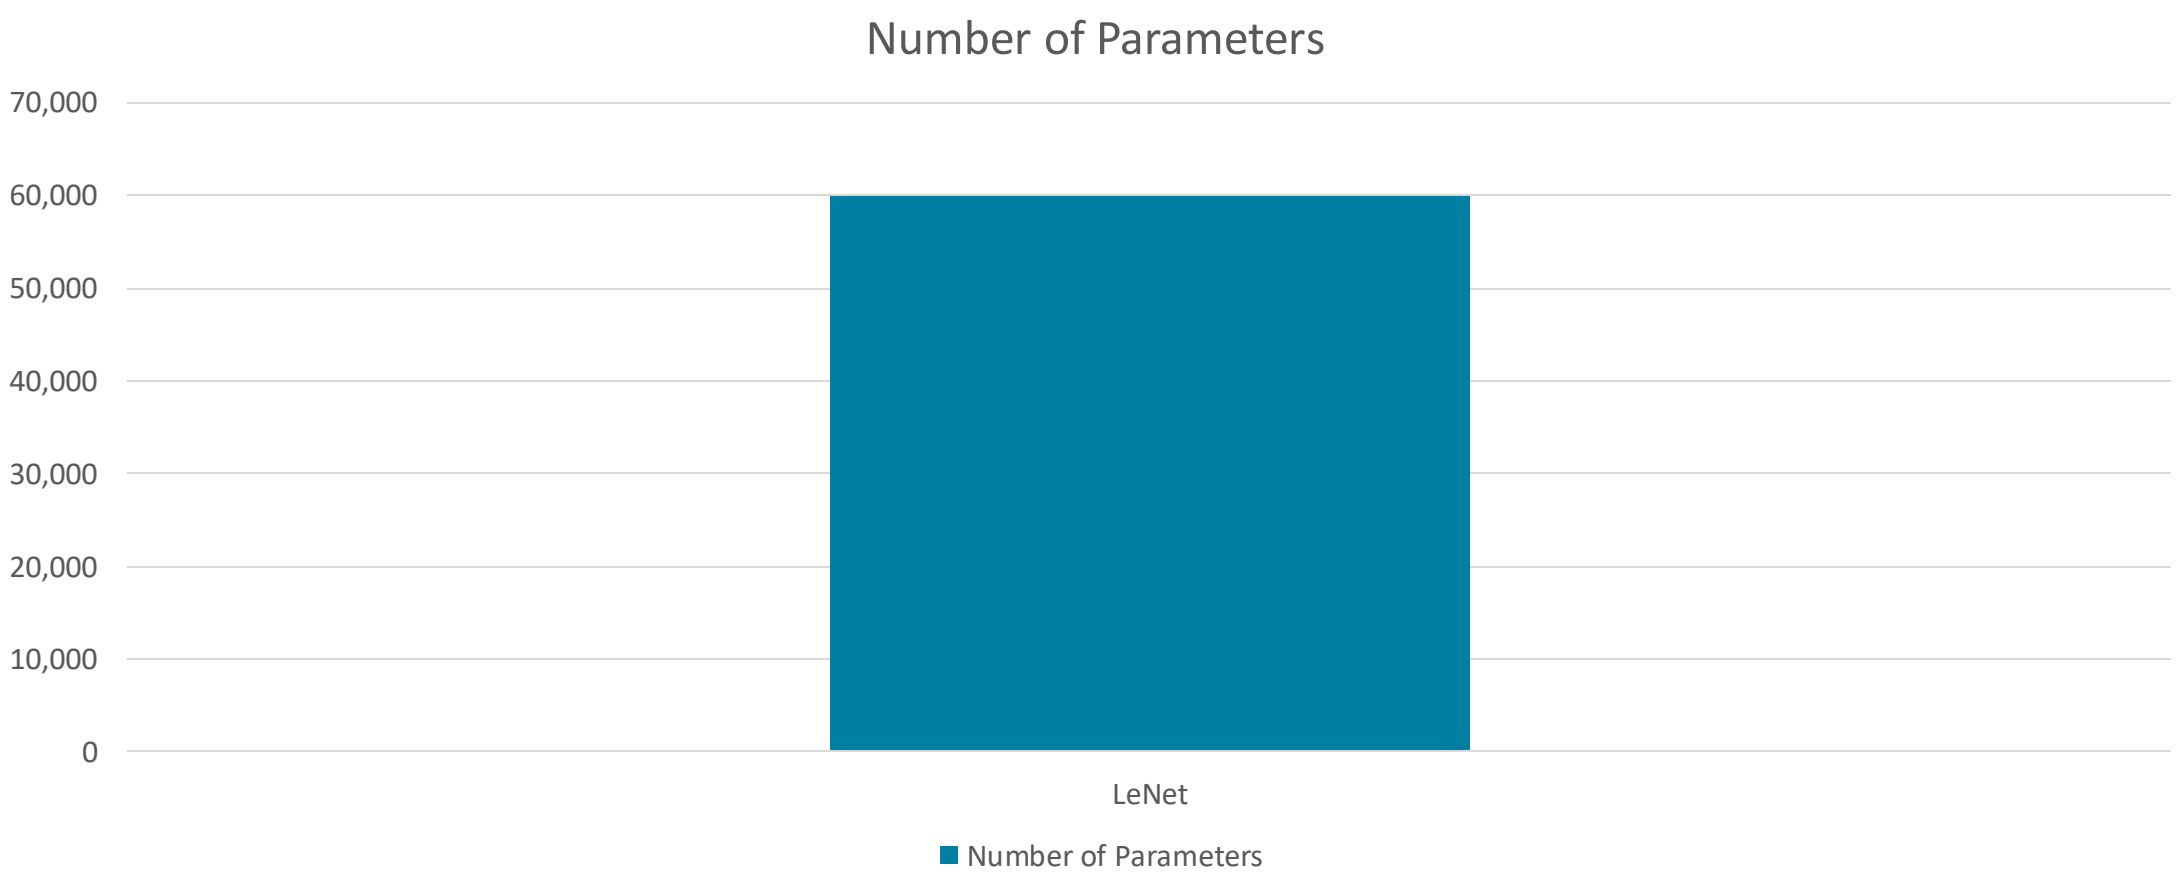
Number of Parameters
70,000
60,000
50,000
40,000
30,000
20,000
10,000
0
LeNet
Number of Parameters

Scale of GPT

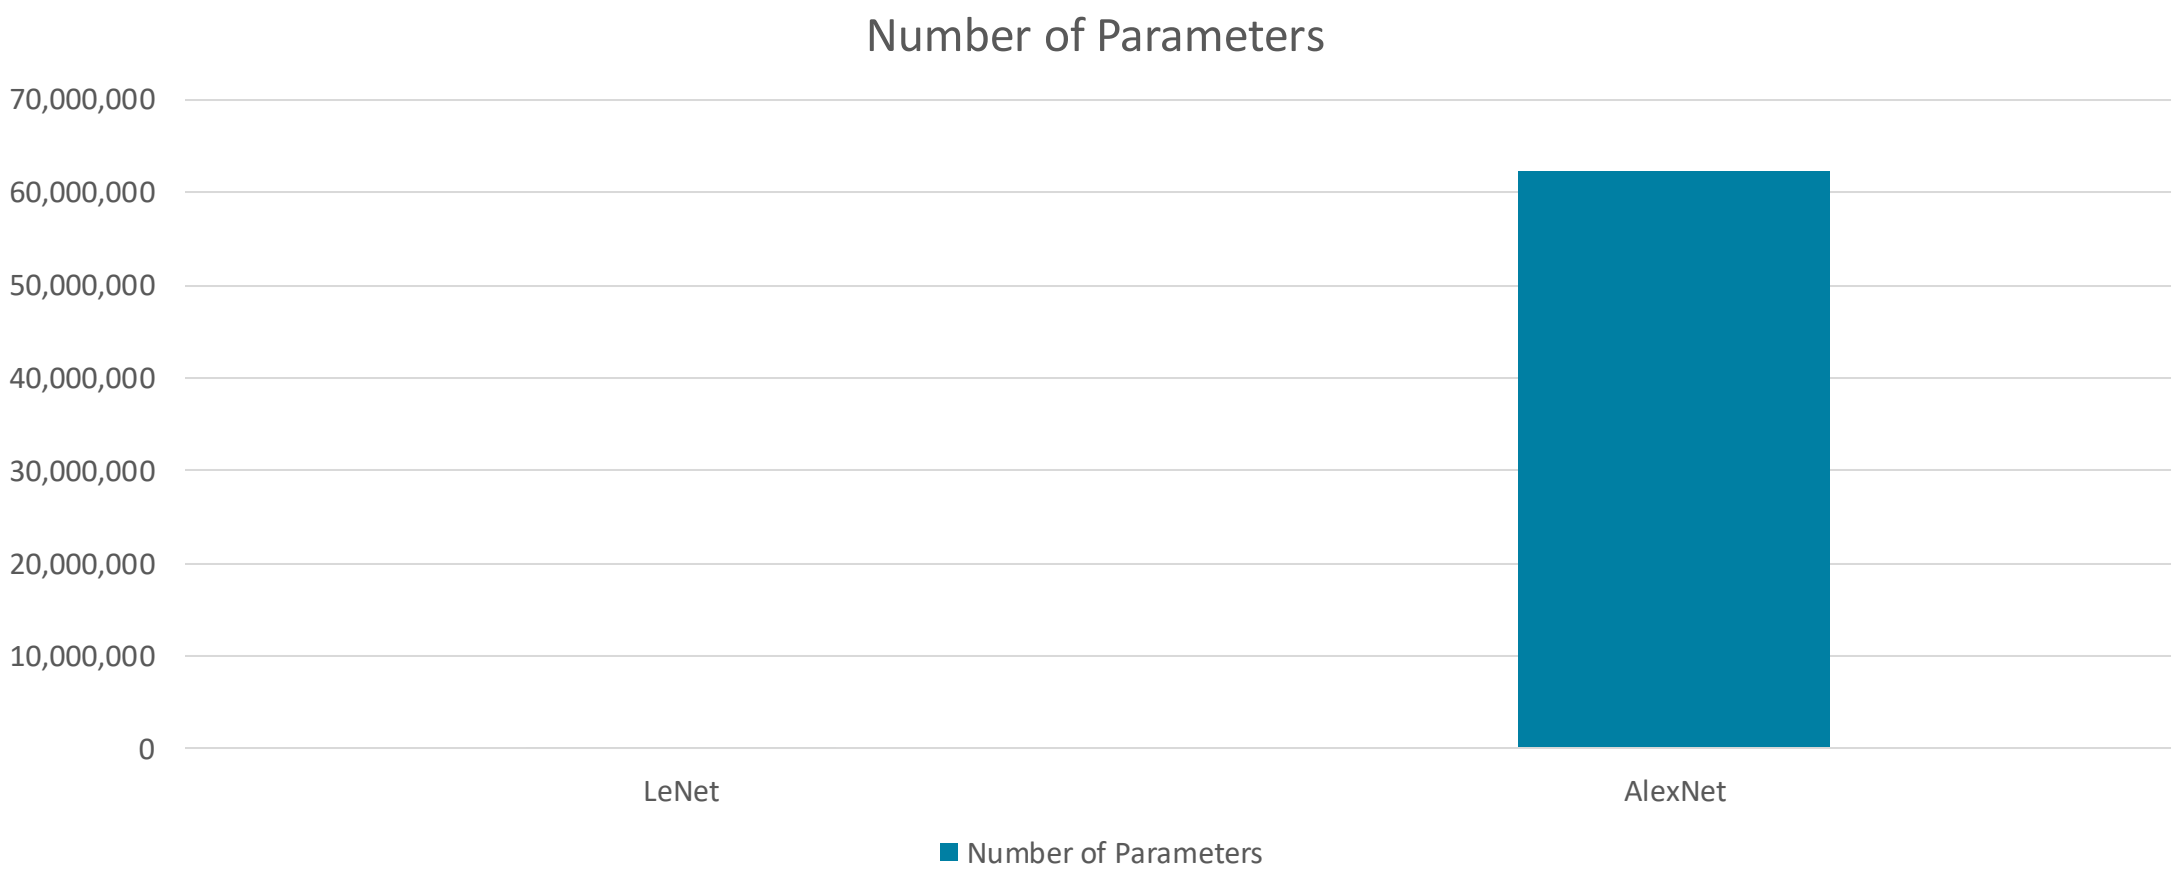
Number of Parameters
70,000,000
60,000,000
50,000,000
40,000,000
30,000,000
20,000,000
10,000,000
0
LeNet
AlexNet
Number of Parameters

Scale of GPT

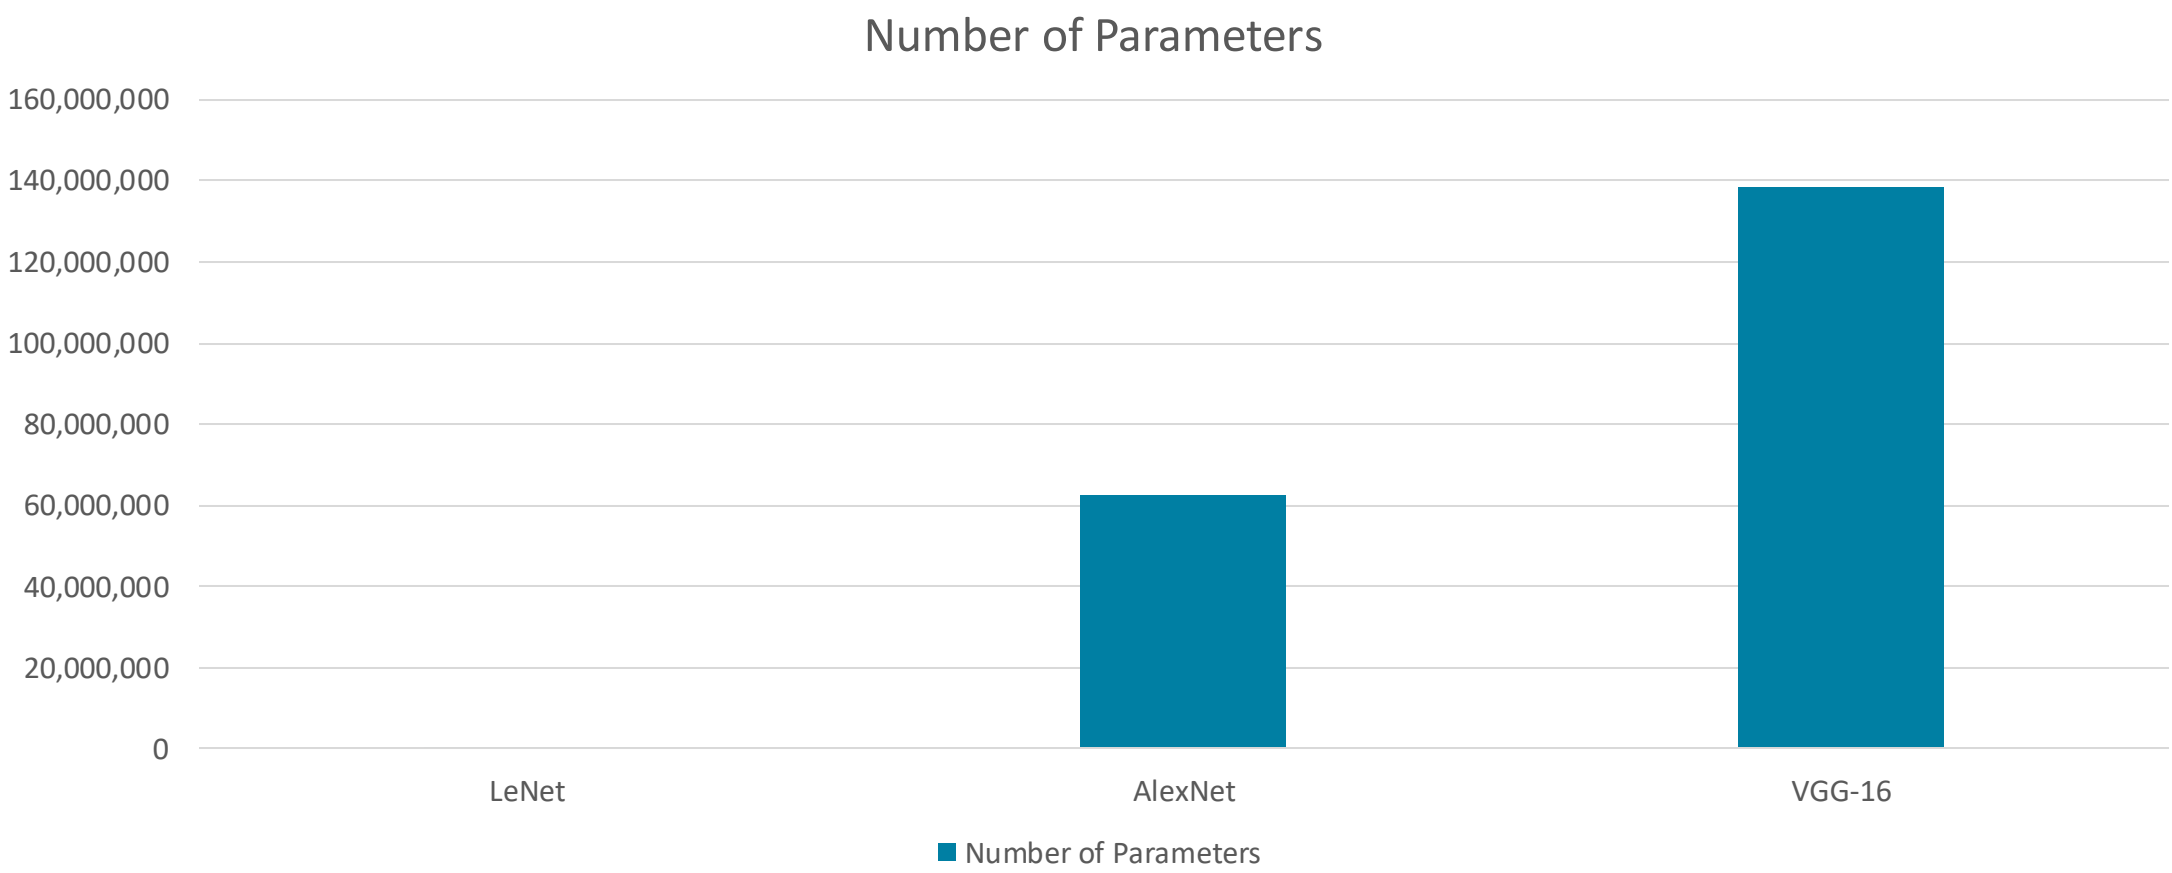
Number of Parameters
160,000,000
140,000,000
120,000,000
100,000,000
80,000,000
60,000,000
40,000,000
20,000,000
0
LeNet
AlexNet
VGG-16
Number of Parameters

Scale of GPT

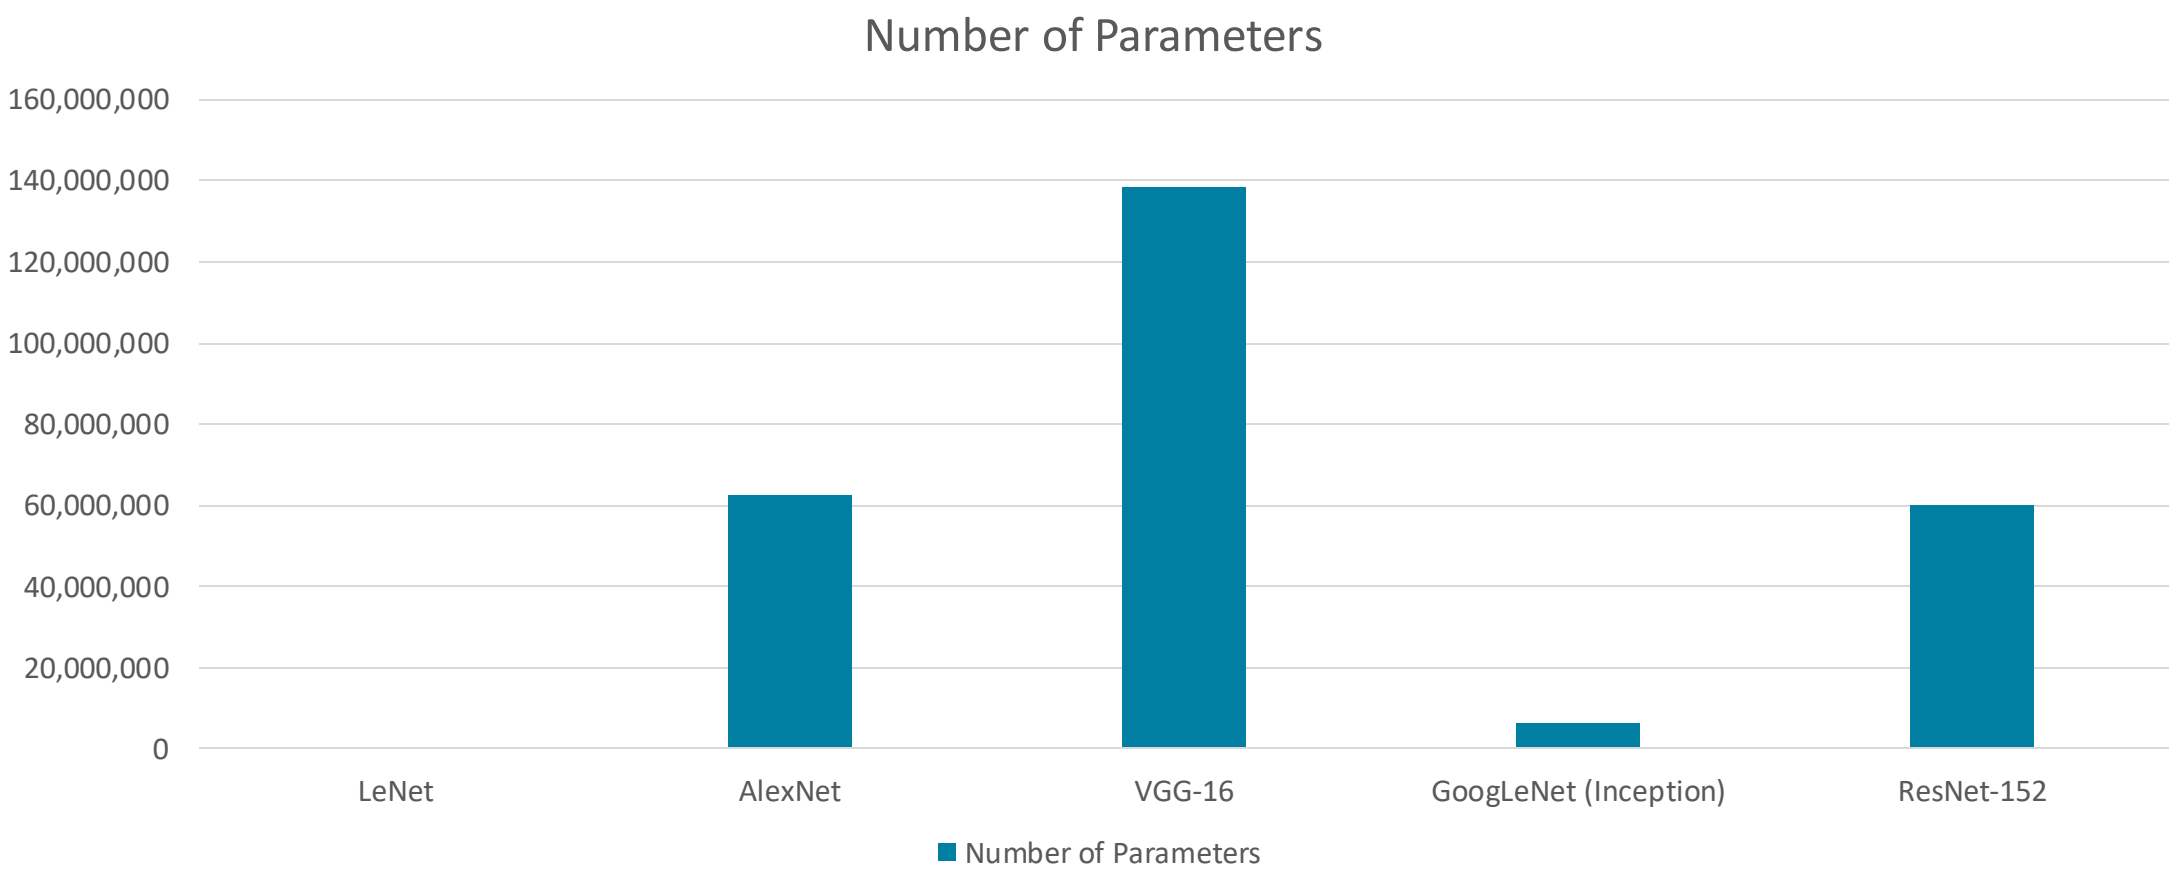
Number of Parameters
160,000,000
140,000,000
120,000,000
100,000,000
80,000,000
60,000,000
40,000,000
20,000,000
0
LeNet
AlexNet
VGG-16
GoogLeNet (Inception)
ResNet-152
Number of Parameters

Scale of GPT

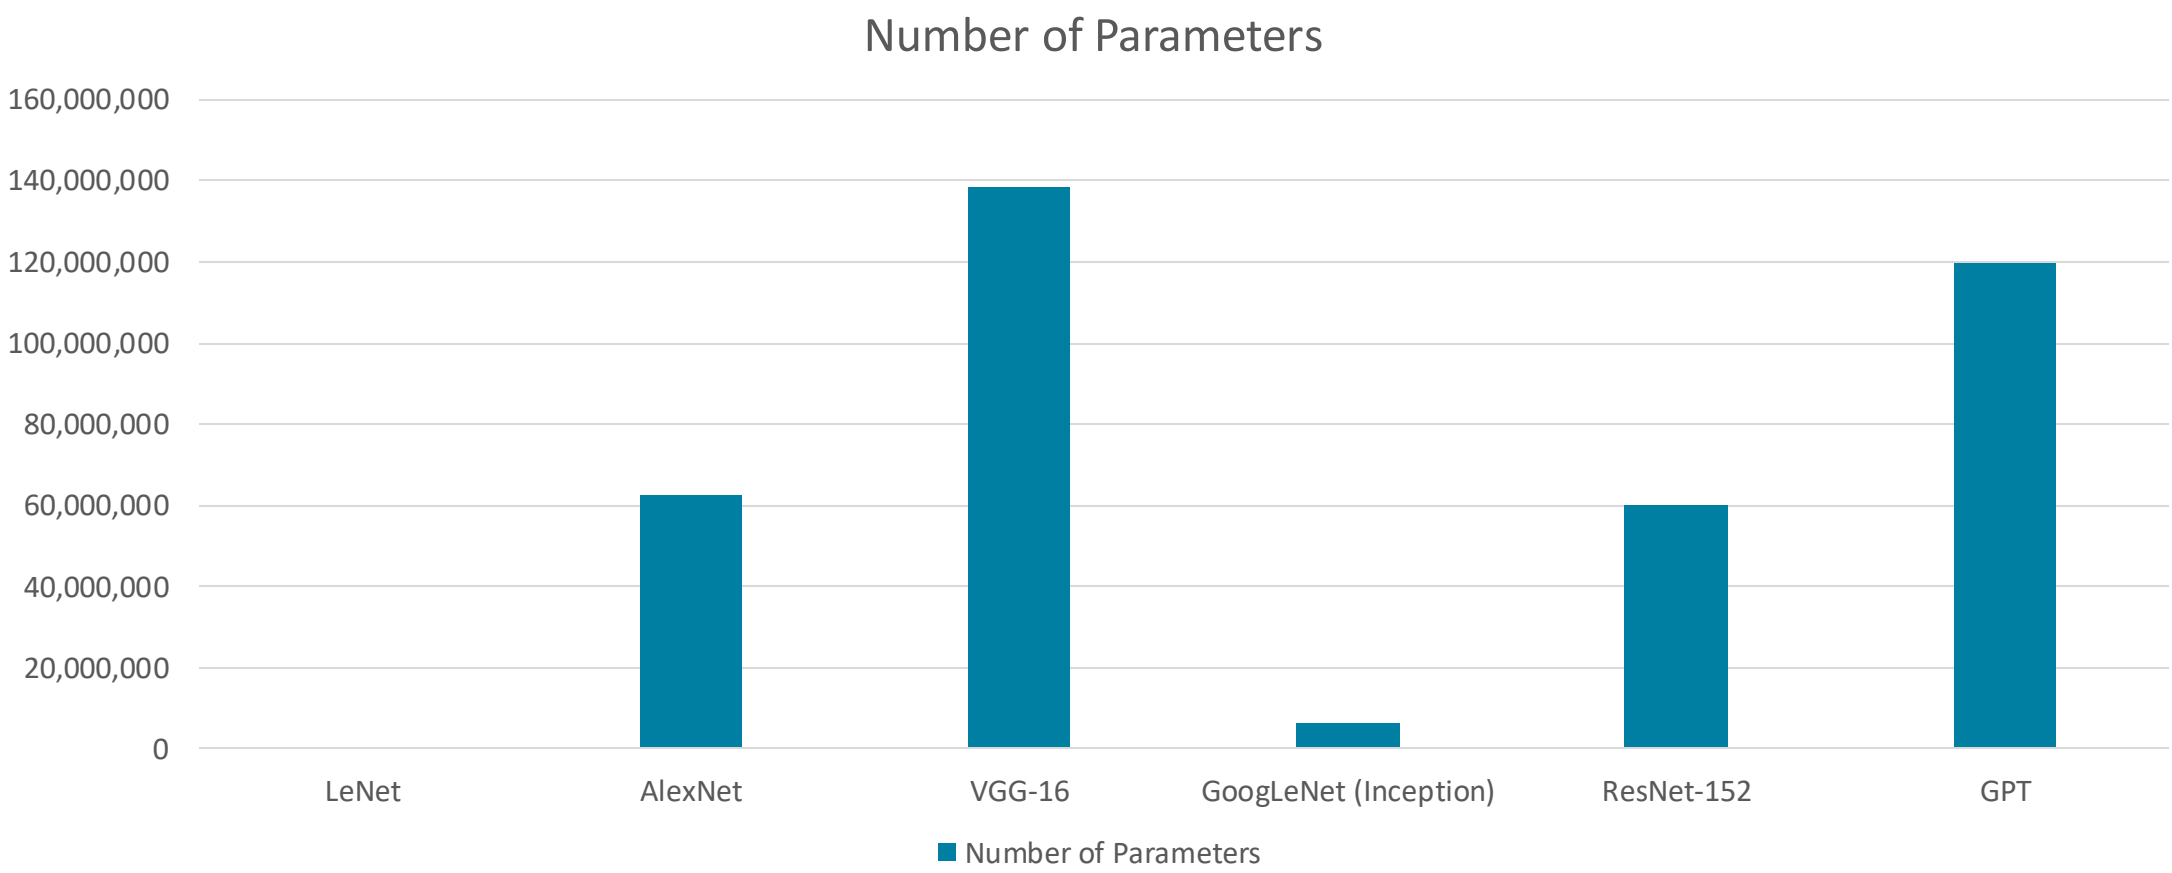
Number of Parameters
160,000,000
140,000,000
120,000,000
100,000,000
80,000,000
60,000,000
40,000,000
20,000,000
0
LeNet
AlexNet
VGG-16
GoogLeNet (Inception)
ResNet-152
GPT
Number of Parameters

Scale of GPT

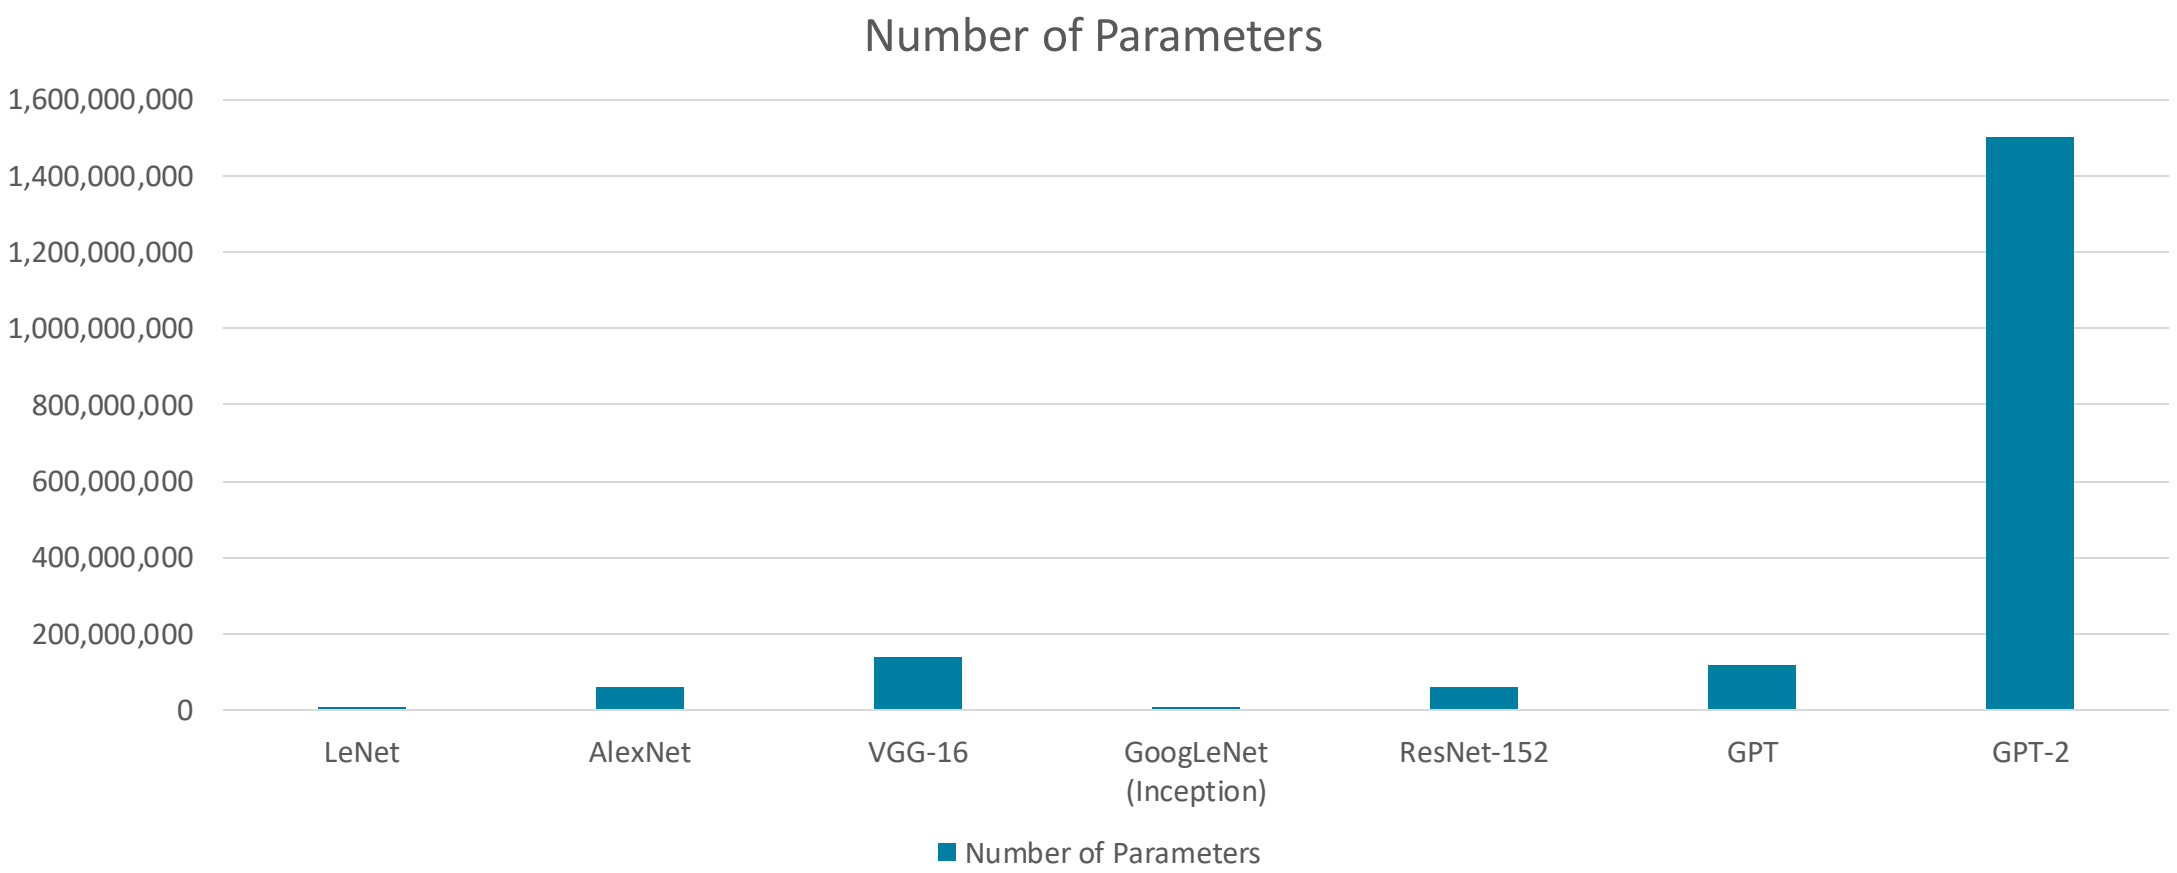
Number of Parameters
1,600,000,000
1,400,000,000
1,200,000,000
1,000,000,000
800,000,000
600,000,000
400,000,000
200,000,000
0
LeNet
AlexNet
VGG-16
GoogLeNet (Inception)
ResNet-152
GPT
GPT-2
Number of Parameters

Scale of GPT

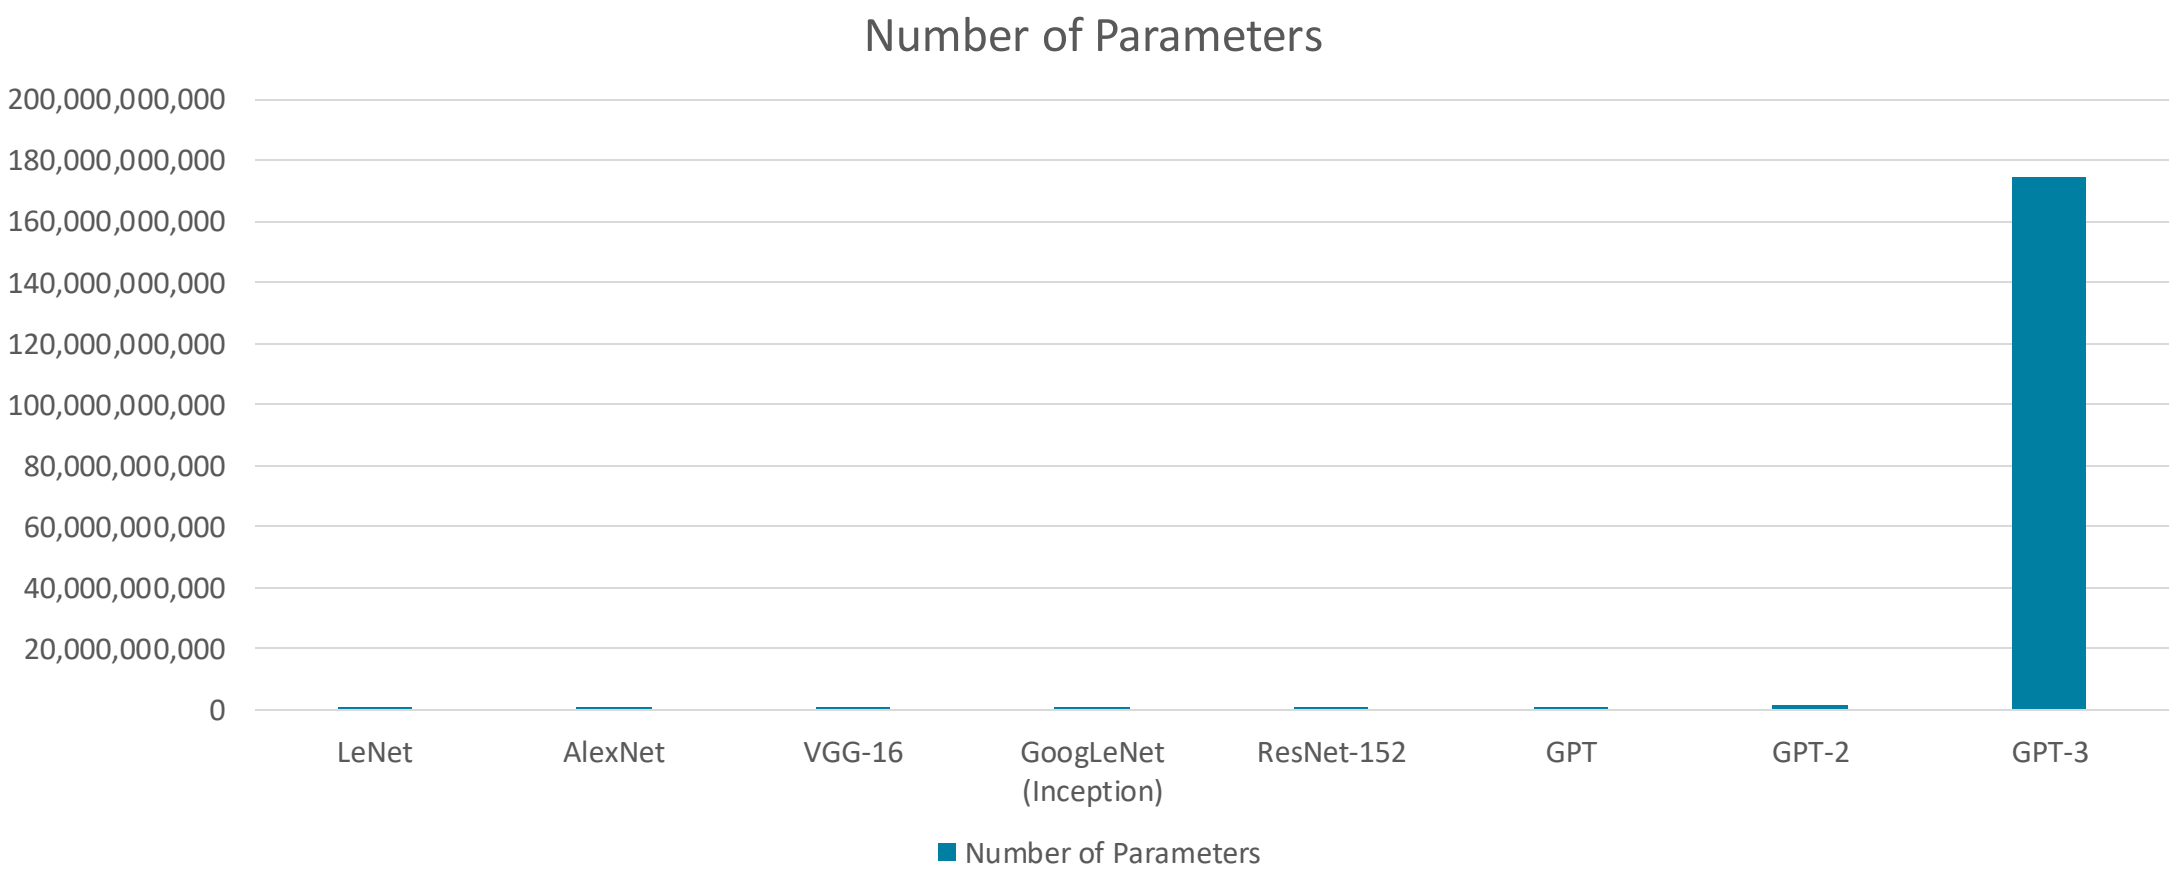
Number of Parameters
200,000,000,000
180,000,000,000
160,000,000,000
140,000,000,000
120,000,000,000
100,000,000,000
80,000,000,000
60,000,000,000
40,000,000,000
20,000,000,000
0
LeNet
AlexNet
VGG-16
GoogLeNet (Inception)
ResNet-152
GPT
GPT-2
GPT-3
Number of Parameters

Scale of GPT

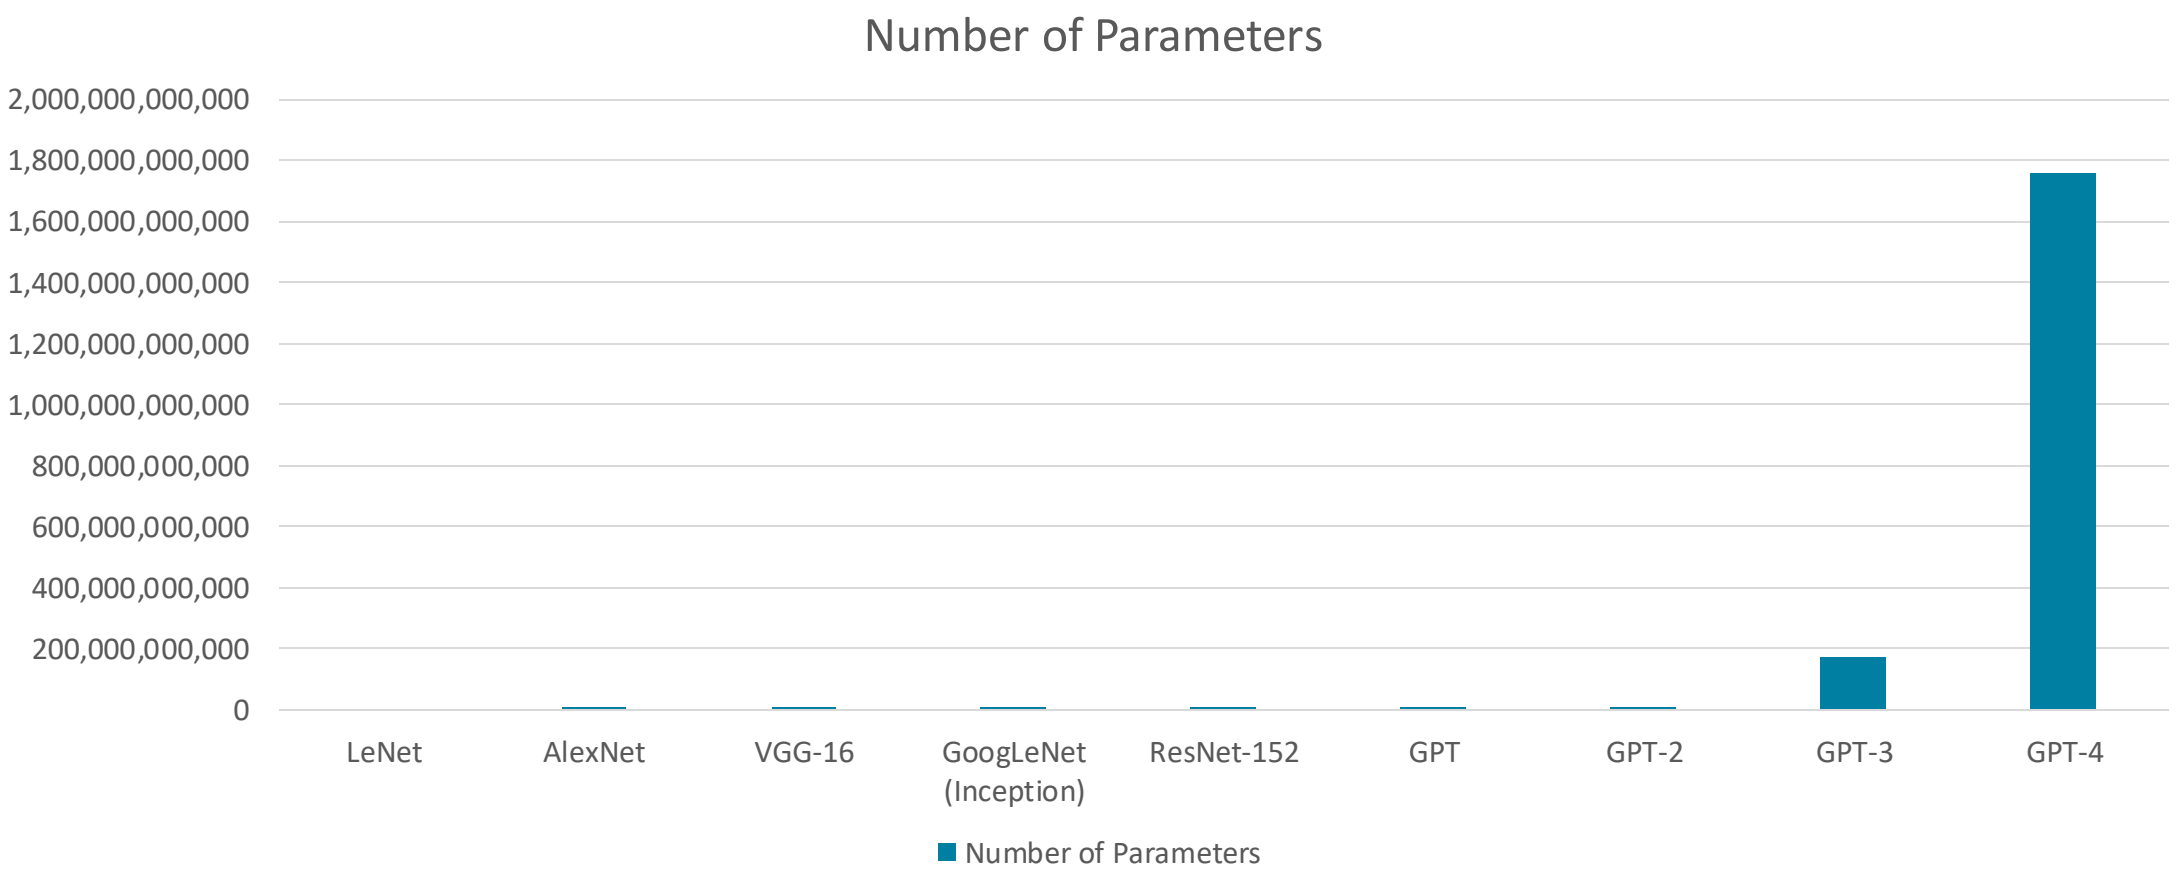
Number of Parameters
2,000,000,000,000
1,800,000,000,000
1,600,000,000,000
1,400,000,000,000
1,200,000,000,000
1,000,000,000,000
800,000,000,000
600,000,000,000
400,000,000,000
200,000,000,000
0
LeNet
AlexNet
VGG-16
GoogLeNet (Inception)
ResNet-152
GPT
GPT-2
GPT-3
GPT-4
Number of Parameters

Scale of GPT

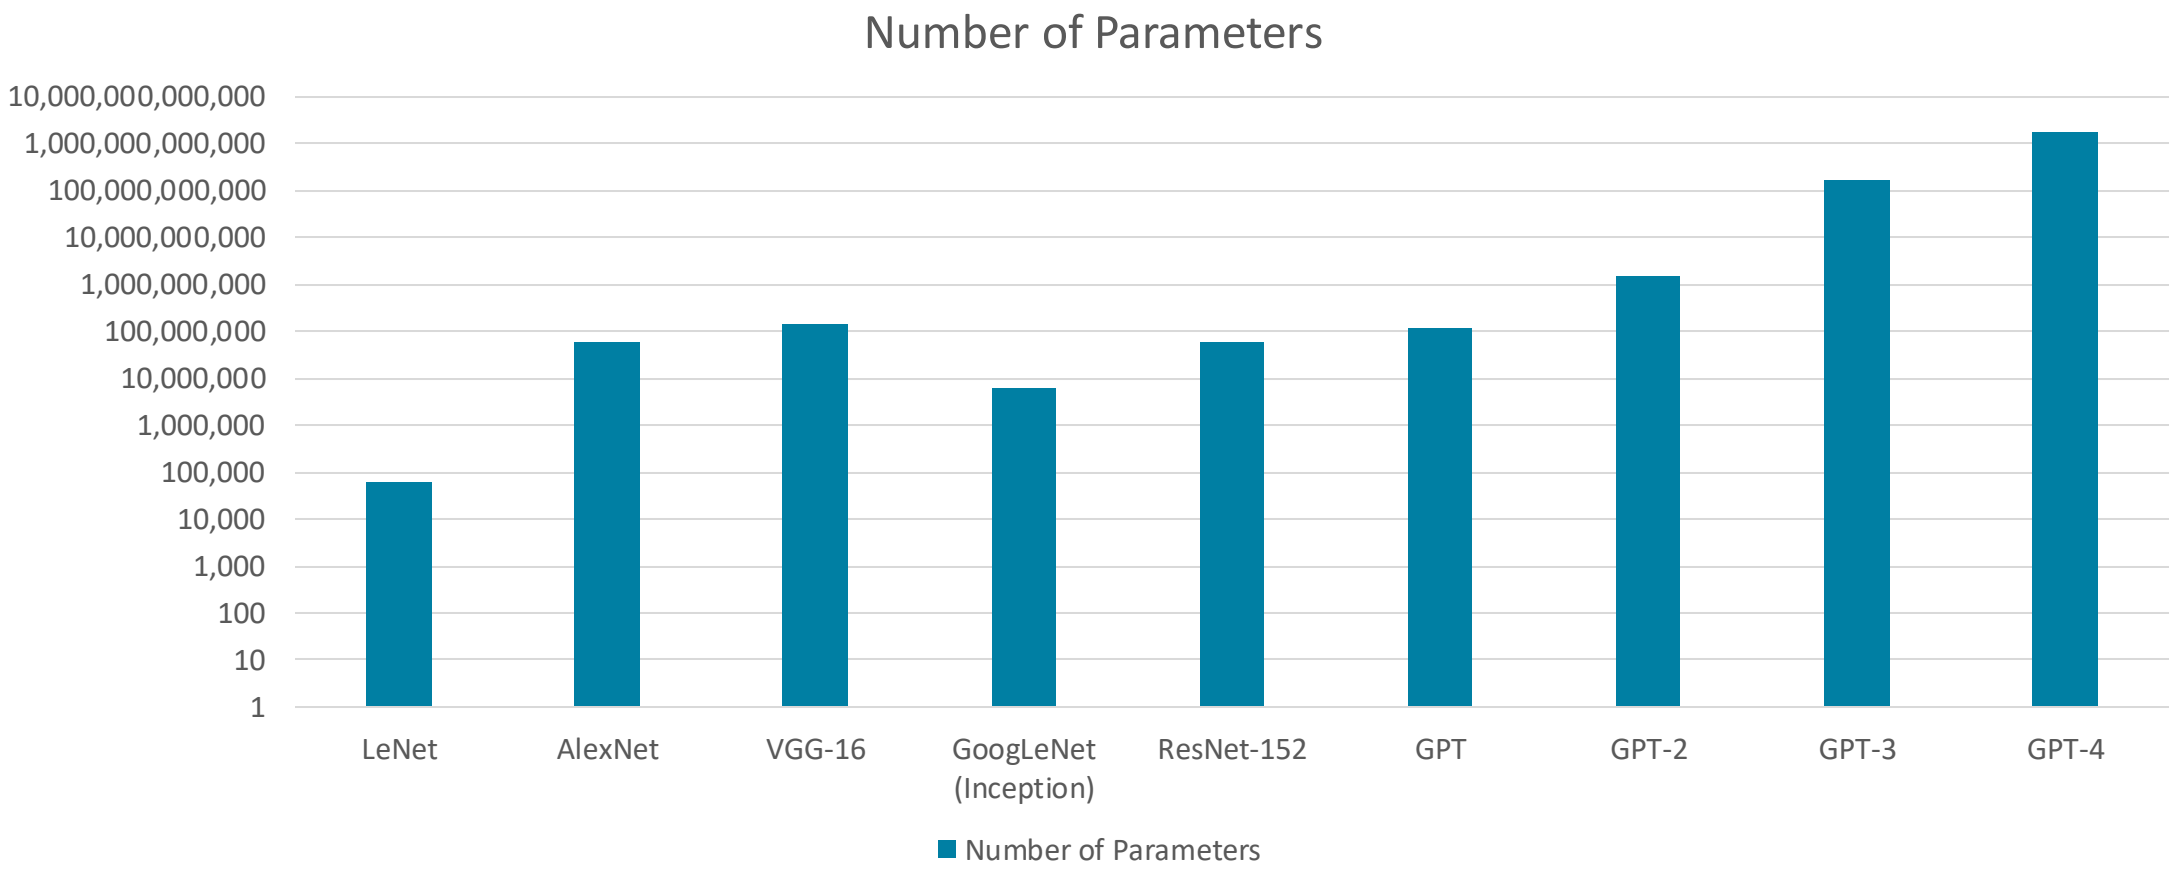
Number of Parameters
10,000,000,000,000
1,000,000,000,000
100,000,000,000
10,000,000,000
1,000,000,000
100,000,000
10,000,000
1,000,000
100,000
10,000
1,000
100
10
1
LeNet
AlexNet
VGG-16
GoogLeNet (Inception)
ResNet-152
GPT
GPT-2
GPT-3
GPT-4
Number of Parameters

GPT's Training Data

- 1 token ≈ ¾ word
- Some datasets are sampled more times than others
- Common Crawl: billions of webpages collected over 7 years
- Webtext2: Dataset of webpages that have been shared on Reddit
- Books1: Free ebooks (?)
- Books2: Secret!
- English Wikipedia

Dataset	Quantity (tokens)	Weight in training mix

The training innovation of ChatGPT

Human annotators write answers to questions

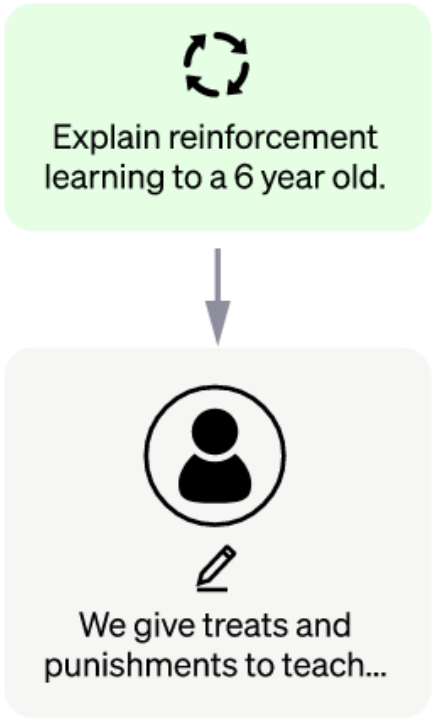

The generalist GPT model is taught from these Q&A pairs

Human annotators write more answers, and someone else ranks them

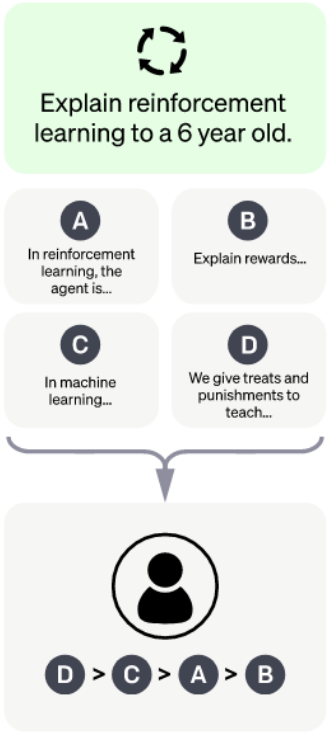

A separate model learns to rate the quality of an answer

GPT writes answers to sampled questions

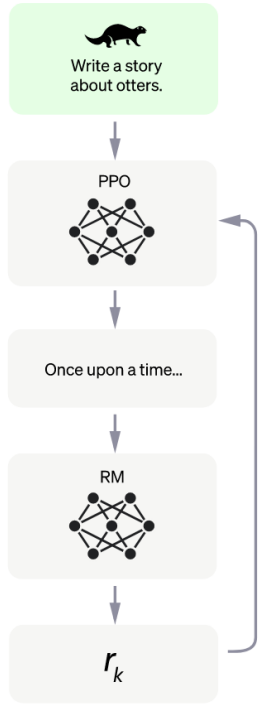

The reward model rates each answer, allowing GPT to keep learning

Figures: https://openai.com/blog/chatgpt/

The training innovation of ChatGPT

No more humans involved!

Human annotators write answers to questions

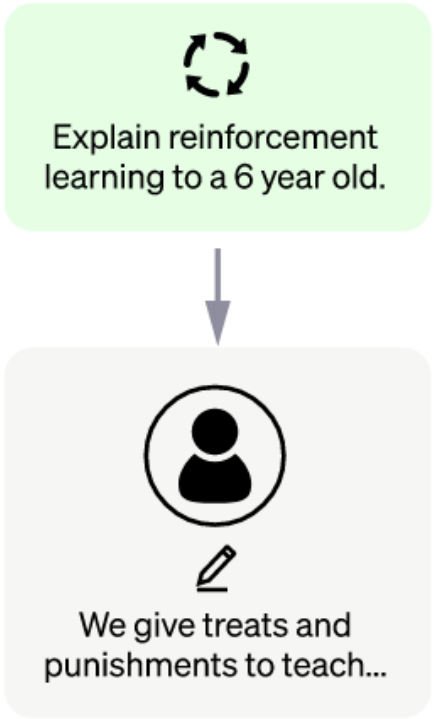

The generalist GPT model is taught from these Q&A pairs

Human annotators write more answers, and someone else ranks them

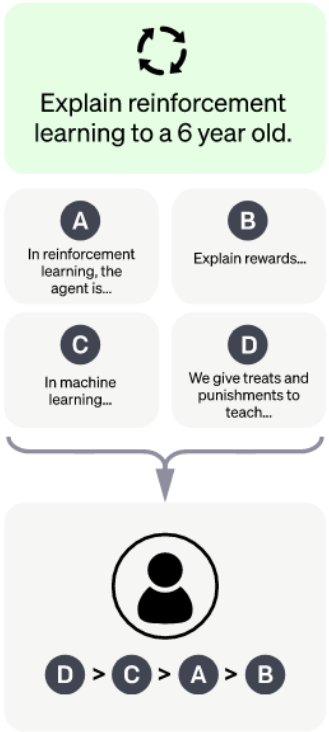

A separate model learns to rate the quality of an answer

GPT writes answers to sampled questions

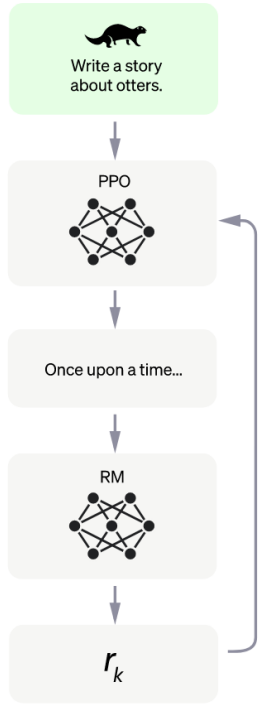

The reward model rates each answer, allowing GPT to keep learning

Figures: https://openai.com/blog/chatgpt/

Winograd Schema

- ”Artificial language processing remains ten years away” – Tom Scott, 2020
- GPT-3 performance: 68.8%
- GPT-4 performance: 94.4%
- Today, 22 models outperform human baselines on the GLUE benchmark